Native Rectangular Linear Assignment with Implicit Rejection: A Jonker–Volgenant Solver without Square Expansion

Ratha SIV

24 October 2026

Abstract

Tracking and matching applications often solve rectangular assignment with a uniform penalty for leaving elements unmatched. Widely used tracking interfaces embed the N × M costs in an (N+M)-square matrix: for 256 rows and 4096 columns, 144.5 MiB of costs for an 8 MiB problem. We describe a Jonker–Volgenant implementation that keeps the real matrix and represents rejection inside its shortest-path search by one terminal distance and one predecessor row. Building on established rectangular and optional-assignment formulations, we derive the one-sided objective and its primal-dual certificate. Rust (JVRUST), hybrid Rust/C++ (JVHYBRID), and C++ (JVCPP) implementations with native batching agree with exhaustive enumeration and are compared with expanded solvers and an exact capped-cost SciPy formulation. Two environments, Python 3.14.9 and 3.11.17 on one Intel/Windows host, each provide 8,946 timing samples across 204 cases, including 5000-square matrices, extreme aspect ratios, structured costs, and rejection sweeps. JVRUST beats capped SciPy in the same nine of ten formulation cases; JVRUST and JVCPP each preserve their ordering against the relevant SciPy baseline in all 201 non-batch cases. Congestion remains a counterexample: capped SciPy is about 9.4 and 9.1 times faster at 1024 × 1280. For small problems, JVRUST's four-worker batching improves on its sequential Python loop by 5.1 and 4.7 times. Repeated memory measurements confirm the cost of square expansion and an orientation-dependent native footprint. These results characterize particular implementations and software environments; they do not isolate a Python-version or programming-language effect.

Keywords: linear assignment; Jonker–Volgenant; rectangular matching; implicit rejection; rejection penalties; research software; Rust; C++; Python

1 Introduction

Every frame, a multi-object tracker must decide which of its tracks continues which new detection. This is a linear assignment problem with two complications that the textbook formulation leaves out: the two sets differ in size, and a pair whose cost exceeds a gating threshold should not be matched at all. An unmatched track may be occluded; an unmatched detection may be a new object. ByteTrack [1] and trackers built on its association code express this with a single call to a Jonker–Volgenant solver from the lap package [2], passing the threshold as a cost limit. The solver then embeds the N × M costs in an (N+M)-square matrix whose dummy rows and columns price each rejection. The same pattern arises wherever correspondences may be missing, in data association, resource allocation, and correspondence estimation. The cost limit prices rejection; it does not require maximizing the number of pairs below the threshold.

The expansion is costly. Consider 256 rows and 4096 columns stored in float64. The original matrix occupies 8 MiB. Zero-padding to 4096 square requires a 128 MiB cost matrix. Modeling unmatched penalties with a 4352 square matrix requires 144.5 MiB. These figures describe the cost matrices alone, before outputs and solver workspace, and the search must also traverse the dummy entries. A native rectangular method can instead operate on the real costs and maintain state proportional to the two set sizes.

Avoiding these transformations is not new in principle. Rectangular assignment, outsourcing, and assignment with insertion and deletion costs already admit compact formulations and specialized augmenting-path algorithms [3–9], and tracking toolboxes accept non-assignment costs directly [10, 11]. This paper

studies an implementation within that established family: a Jonker–Volgenant solver whose rejection option needs only constant-size search state, which accepts the threshold arguments of the `lap.lapjv` interface and solves batches natively. We evaluate it against both expanded interfaces and an exact capped-cost SciPy formulation on the original matrix dimensions.

This paper describes such a solver. It combines a Jonker–Volgenant square strategy, derived from the scalar src-d/lapjv approach [12], with native rectangular matching and an implicit rejection terminal that represents a uniform unmatched penalty without dummy rows or columns. The solver is implemented three times with identical semantics: entirely in Rust (JVRUST), in Rust with a C++ square-completion kernel (JVHYBRID), and entirely in C++ (JVCPP). JVRUST is distributed as the laprust package [13] and as the `lapjvr` family of LAPX [14]; JVHYBRID is LAPX's `lapjvh` family; JVCPP is the lapcpp package [15]. All three expose the same single-matrix and batch interfaces.

The paper makes four contributions.

- **Derivation.** A self-contained derivation of the uniform-penalty objective in one-sided form and of the primal-dual certificate the solver satisfies, related to established optional-assignment formulations (Section 3).
- **Implementation.** A Jonker–Volgenant solver that combines rectangular reductions with a rejection terminal stored as one distance and one predecessor row, linear solver state on suitable wide input, the threshold arguments of `lap.lapjv`, and native batch workers (Section 4).
- **Baseline evaluation.** Comparison with an exact capped-cost formulation on the original dimensions, a specialization of established outsourcing and error-correction reductions [7, 8], alongside expanded formulations (Sections 3.5 and 6.5).
- **Empirical characterization.** Rust, hybrid Rust/C++, and C++ implementations measured in two Python environments on one host, covering batching, repeated memory trials, large and extreme shapes, structured costs, rejection sweeps, and workloads where SciPy is faster (Sections 6 and 7).

Section 2 positions the solver among existing methods and interfaces. Section 3 formulates the problem and proves the reformulations, and Section 4 describes the implementation. Sections 5 and 6 present the experimental design and results, Section 7 discusses what they imply for practitioners, and Section 8 lists the limitations.

2 Background and related work

2.1 Assignment algorithms and implementation lineage

Jonker and Volgenant's algorithm combines initialization, reductions, and a specialized shortest augmenting path procedure [3]. Initial stages often construct much of the matching cheaply. Remaining free rows require augmenting paths that can reassign existing pairs. A square matrix therefore does not eliminate the augmentation phase; it only removes the need for a dimensional transformation.

We use completion to mean the augmenting-path phase that finishes a partial matching. It does not refer to padding the cost matrix.

The baselines are established Jonker–Volgenant families, evaluated through their LAPX implementations [14]. JVX (`lapjv`, `lapjvx`, `lapjvxa`) uses the gatagat/lap-derived dense core [2]; this study times `lapjvxa`, which returns assignment pairs. JVS (`lapjvs`, `lapjvsa`) uses the src-d/lapjv-derived implementation [12]. JVC (`lapjvc`) adapts the classic dense shortest-path implementation in py-lapsolver [16]. These implementations have different initialization and numeric details. Their shared use of augmenting paths does not imply identical kernels.

SciPy provides a native rectangular linear assignment solver. Its documentation describes a modified Jonker–Volgenant algorithm without initialization and references Crouse's treatment of rectangular assignment

[4, 17]. It is therefore an important counterpoint to the padded interfaces: rectangular solving without square padding already exists independently of the proposed solver. Volgenant gives an earlier adaptation of LAPJV to non-square problems [9]. Bijsterbosch and Volgenant refine it and develop specialized rectangular implementations for k-cardinality and outsourcing problems, including a comparison with a capped-cost formulation that keeps the original dimensions [8]. That unexpanded outsourcing formulation originates with Vander Wiel and Sahinidis, whose shortest augmenting path algorithm treats external assignment through a single unrestricted arc [18]. Their outsourcing search is a direct predecessor of the optional-assignment search studied here, not merely an alternative problem definition.

2.2 Optional assignment and compact representations

Optional assignment also has clear precedents. Bougleux and Brun formulate linear sum assignment with edition using a null element on each side rather than a full square expansion [5]. Their augmenting paths can terminate at a null element, representing removal or insertion. The associated LSAPE toolbox supplies C++ implementations, including null-terminal path routines, in releases dated 2017 [19].

Jones, Chawdhary, and King adapt Volgenant–Jonker methods to graph-edit matching [6]. Their compact representation avoids explicitly storing the redundant zero and infinite regions of the expanded matrix, using only (n+1)(m+1)-1 entries. Bougleux and colleagues subsequently describe a compact reduction from assignment with error correction to a rectangular assignment problem without cost constraints of the earlier reductions [7].

The present solver is an implementation specialization within this family. It retains the real cost matrix and represents rejection during augmentation by one terminal distance and one predecessor row; Section 4.3 gives the update. We do not claim a new optional-assignment objective, a new null-terminal principle, or the first Jonker–Volgenant implementation that avoids storing a square cost matrix. The code descends from the JVS lineage; conceptual predecessors are cited independently of code provenance.

Application interfaces also accept non-assignment costs. MATLAB's `assignjv`, `assignmunkres`, and `assignauction` take a scalar cost of non-assignment charged per unmatched endpoint [11]. The Tracker Component Library includes a Jonker–Volgenant routine for missed detections that never builds the augmented matrix but keeps search state for one virtual row per target [10]. The implementation combines a constant-size rejection terminal with rectangular reductions, the threshold arguments of `lap.lapjv`, and native batch execution.

The compact reduction in [7] is also an executable comparison point. For uniform endpoint penalties, its cost factorization reduces, up to a constant shift, to capping real costs at the combined rejection price. Section 3.5 gives a self-contained argument for this specialization, including finite signed costs, and Section 6.5 evaluates it with SciPy. This provides a baseline without square expansion; it does not reproduce every optimization in the authors' LSAPE software.

2.3 Public interfaces and transformed problems

Table 1 summarizes how the evaluated dense solvers handle rectangular inputs and unmatched penalties. With no finite limit, an $N \times M$ required problem assigns every element on the smaller side, assuming feasibility. A finite cost limit changes the objective and may reduce the number of pairs.

The sparse LAPMOD interface of LAPX is outside the performance study because it accepts a different representation and requires square sparse problems.

Table 1: Native interface behavior of the evaluated solvers. In the primary finite-limit suite, JVS, JVC, and SciPy receive explicitly constructed square-augmented problems. Sections 6.5 and 6.6 additionally evaluate SciPy with dummy-column and original-dimension capped-cost formulations. The wrappers preserve the objective without adding a native rejection option to the underlying required-assignment API.

Solver	Required rectangular input	Finite unmatched penalty
JVX	Zero-padding to max(N,M) square	Internal (N+M) square augmentation
JVS	Wrapper zero-padding to max(N,M) square	No public limit parameter
JVC	Wrapper zero-padding to max(N,M) square	No public limit parameter
Proposed solver (three implementations)	Native rectangular solving	Implicit rejection terminal
SciPy	Native rectangular solving	No equivalent public limit parameter

3 Problem formulation

3.1 Required rectangular assignment

Let C be an N × M real cost matrix, with forbidden edges excluded from the admissible edge set E. Orient the matrix so that r = min(N,M) and c = max(N,M), with r ≤ c. In (1) and (4), $x_{ij} = 0$ for $(i, j) \notin E$. A required rectangular assignment solves

$$\min_{x} \sum_{(i,j)\in E} C_{ij}x_{ij}, \quad \sum_{j} x_{ij} = 1, \quad \sum_{i} x_{ij} \le 1, \quad x_{ij} \in \{0, 1\}. \tag{1}$$

The solver must match all r rows if this is feasible. Surplus columns remain unused. If the original matrix is tall, the returned indices are mapped back to the original orientation.

3.2 Finite-limit assignment

Let L be a finite limit and A a matching with K pairs. The finite-limit objective is

$$F(A) = \sum_{(i,j)\in A} C_{ij} + \frac{L}{2}(r + c - 2K). \tag{2}$$

Each unmatched vertex incurs L/2. In the square augmentation used by JVX, the real costs occupy the upper-left block. Real-to-dummy and dummy-to-real blocks cost L/2, while the dummy-to-dummy block is zero.

Equation (2) is not equivalent to solving a required assignment and then discarding expensive pairs. For example, let

$$C = \begin{pmatrix} 0 & 0.9 \\ 0.9 & +\infty \end{pmatrix}, \qquad L = 1.$$

Two real pairs cost 1.8. Selecting only the zero-cost pair leaves two unmatched vertices and costs 1.0 under (2). Both nonzero finite edges are below the threshold, yet the smaller matching is preferable.

At equality, a real edge of cost L can tie with rejecting both endpoints. The objective does not impose maximum cardinality among equal-cost optima. Applications that require a particular inclusive-threshold or cardinality rule must specify it separately.

3.3 One-sided rejection equivalence

Proposition 1. In exact arithmetic, minimizing (2) is equivalent to minimizing

$$P(A) = \sum_{(i,j)\in A} C_{ij} + L(r - K) \tag{3}$$

on the oriented real rectangular problem.

Proof. Rearranging the rejection term gives

$$F(A) = \sum_{(i,j)\in A} C_{ij} + L(r - K) + \frac{L}{2}(c - r).$$

The final term is independent of A. Consequently, F and P have the same minimizers. The argument holds for any finite real L, including zero and negative values, and includes empty dimensions. ■

The transformation is useful because every smaller-side row can either select a real column or take a rejection option costing L. No separate dummy column is needed for each possible rejection. The constant must be restored when comparing the full penalized objective, but it does not affect the optimizer.

The public returned total is different again: it sums only selected original input costs and excludes rejection penalties. Our validation and benchmark checks reconstruct the penalized objective rather than confusing it with this reported total.

3.4 Primal-dual certificate

For an effective row penalty p, introduce a rejection variable z_i:

$$\begin{aligned} \min_{x,z} \quad & \sum_{(i,j)\in E} C_{ij}x_{ij} + p\sum_i z_i \\ \text{subject to} \quad & \sum_j x_{ij} + z_i = 1, \qquad \sum_i x_{ij} \le 1, \qquad x_{ij}, z_i \ge 0. \end{aligned} \tag{4}$$

This is a network-flow formulation: each row contributes one unit, real columns have unit capacity, and a row may send its unit directly to the terminal at cost p. The integral capacities yield integral optimal solutions. Its dual is

$$\max_{u,v} \sum_i u_i + \sum_j v_j, \quad u_i + v_j \le C_{ij} \quad ((i,j) \in E), \quad u_i \le p, \quad v_j \le 0. \tag{5}$$

Proposition 2. A feasible matching/rejection solution is optimal if a feasible dual satisfies tightness on matched real edges, $u_i = p$ for rejected rows, and $v_j = 0$ for unused columns.

Proof. The primal-dual gap can be written as

$$P - D = \sum_{(i,j)\in E} (C_{ij} - u_i - v_j)x_{ij} + \sum_i (p - u_i)z_i + \sum_j v_j \left(\sum_i x_{ij} - 1 \right). \tag{6}$$

Every term is nonnegative by primal and dual feasibility. The stated conditions make each term zero, so the primal and dual objectives coincide. Weak duality establishes optimality. ■

For required rectangular assignment, rejection variables and the constraints $u_i \le p$ are omitted. The non-positive column potentials and zero potentials on unused columns remain important. Ordinary square reductions cannot be transferred blindly to a rectangular implementation.

These propositions establish the formulation and a final optimality certificate. They are not a formal verification of every implementation transition or a floating-point error bound.

3.5 Equivalent formulations for a stronger baseline

The one-sided objective also permits two formulations for an ordinary required rectangular solver. First, append r dummy columns of cost L to the oriented real matrix. The resulting r × (c+r) problem assigns every row either to a real column or to one of the distinct dummy columns. Removing dummy pairs recovers an optimal finite-limit matching under (3).

Second, no additional columns are necessary when rejection costs are uniform. Define a capped matrix on the original r × c dimensions:

$$\widehat{C}_{ij} = \min(C_{ij}, L), \qquad \min(+\infty, L) = L. \tag{7}$$

Proposition 3. Let B be a minimum-cost required matching of all r rows under (7). Retain only its pairs whose original cost is strictly below L. The retained matching minimizes (3); rejecting equality ties is one valid optimality convention.

Proof. For any required matching B, let A be its retained subset. Every discarded pair costs L in the capped problem, so the capped cost of B is exactly P(A). Conversely, any finite-limit matching A can be extended to r pairs because c≥r. Each of the added pairs costs at most L after capping, and each original pair costs at most its original cost. Thus some required matching has capped cost at most P(A). Taking minima in both directions proves equality of the optimal values and optimality of the retained subset. The argument holds for finite real L and finite admissible costs, including negative values; forbidden original pairs are discarded. ■

This is the uniform-penalty case of the substitution factorization underlying [7], stated here to make the comparison reproducible. Capping occurs before optimization. It is therefore different from solving the original required problem and then discarding expensive pairs, which fails in the example in Section 3.2. The capped and dummy-column baselines are equivalent in exact arithmetic; their floating-point behavior is assessed separately by the validation experiments.

4 Implementation

4.1 Shared architecture and three implementations

Each implementation consists of a thin Python extension and a native core. The extension validates shape, real numeric type, precision, and limit options, prepares a contiguous working representation where necessary, and invokes the core. The single-matrix and batch interfaces, return contracts, and error behavior are identical across the three implementations. Table 2 lists them.

Table 2: Implementations of the proposed solver. The main core covers initialization, reductions, and the rectangular and rejection search. All three implement the same algorithm and objective. JVRUST is distributed in two packages with identical native code; the experiments time laprust and validate both.

Label	Main core	Square completion	Distribution
JVRUST	Rust	Rust	laprust 1.0.0 [13]; LAPX 0.11.1 `lapjvr` [14]
JVHYBRID	Rust	C++	LAPX 0.11.1 `lapjvh` [14]
JVCPP	C++	C++	lapcpp 1.0.0 [15]

Interface compatibility here means that the single-matrix entry points accept the `extend_cost`, `cost_limit`, and `return_cost` arguments of `lap.lapjv` with the same objective: the finite limit prices unmatched endpoints, and the reported total excludes those penalties. They return matched row and column index arrays, or pairs, rather than lap's full-length per-row and per-column assignment vectors, so callers of lap need a small output adaptation. Compatibility does not imply identical assignments when optima tie, or identical behavior across precision modes.

The native dispatcher selects the square kernel only for square inputs with an infinite limit. The square kernel's initialization and reductions are followed by an augmentation phase that completes the matching. JVRUST executes this phase in Rust. JVHYBRID calls a C++ completion kernel that operates on Rust-owned workspace and preserves the completion strategy and operation order. JVCPP implements the whole algorithm in C++, including the same C++ completion routine. This design permits two comparisons: of completion kernels within otherwise identical Rust code (JVRUST against JVHYBRID), and of complete Rust and C++ implementations of one algorithm (JVRUST against JVCPP).

All rectangular inputs and finite-limit inputs, including square matrices with a finite limit, use the rectangular core. JVRUST and JVHYBRID share that Rust core, so differences between them on these paths cannot be interpreted as evidence of different rectangular algorithms. JVCPP executes the same rectangular algorithm in C++.

4.2 Native rectangular reductions and search

The rectangular core stores row and column mappings separately, together with potentials, predecessors, distances, and free-row lists. Column potentials are constrained to be non-positive, with zero potentials on unused columns. Bounded row reductions lower a column potential only while assigning that column, preserving the relevant occupancy condition.

After reductions, each remaining free root is introduced with a feasible row potential obtained from its available reduced costs and, in finite-limit mode, the rejection penalty. The root's reduced costs are retained as the initial shortest-path distances, so the first relaxation does not read the cost row again. A shortest-path search then scans real columns. Reaching a free column completes an augmenting path. Reaching an occupied column continues through its assigned row, allowing previous pairs to be rearranged.

The implementation has specialized reduction policies for square finite-limit inputs and sufficiently large near-square required inputs, and grouped scans for some minima. These choices are engineering policies. Their individual benefits and sensitivity are not isolated by the present public-interface experiments.

4.3 Implicit rejection terminal

For finite-limit matching, the search compares real-column paths with a rejection terminal. A real edge has reduced cost $C_{ij} - u_i - v_j$. Leaving a row unmatched has reduced cost $p - u_i$. The implementation stores the best terminal distance and its predecessor row using the variables `dummy_distance` and `dummy_row`.

More precisely, let δ_i be the shortest-path distance at which row i is reached; the root has $\delta = 0$. When row i is reached, the rejection candidate $\delta_i + p - u_i$ updates the terminal distance $D_\perp$, the smallest such value so far, and its row. Each unsettled real column j is relaxed to $\min(d_j,\ \delta_i + C_{ij} - u_i - v_j)$. If $D_\perp$ is no greater than the smallest unsettled column distance, rejection is selected; otherwise that column is settled, and the search either ends at a free column or continues through the column's matched row. The terminal therefore adds one distance and one predecessor row to the ordinary search state.

Let D be the distance of the selected endpoint. The root's potential increases by D. Every settled column j has its potential lowered by $D - d_j$ and, if matched, its row's potential raised by the same amount. Assignments are then reversed from the free column, or from the terminal's predecessor row, back to the root. A free column ending the path has $d_j = D$ and keeps its zero potential, so column potentials stay non-positive and unused columns stay at zero.

If rejection is selected, the number of real pairs need not increase: the path can replace one row's real assignment and leave another row unmatched.

Algorithm 1 summarizes the distinction between this representation and explicit matrix augmentation.

```
Orient C so that r <= c.
Initialize mappings, potentials, and free-row state.
Apply bounded row reductions, including rejection choices in finite-limit mode.
For every row still requiring processing:
    Establish the root row potential from real reduced costs and finite p.
    Search shortest alternating paths over real columns.
    In finite-limit mode, also track the best rejection terminal.
    Stop at a free real column or the selected rejection terminal.
    Update potentials and reverse assignments along the path.
Map real pairs back to the original matrix orientation.
```

Algorithm 1: High-level native rectangular control flow with implicit rejection. The implementation additionally specifies reduction budgets, finite-value checks, tie handling, and workspace reuse.

4.4 Storage, batching, and operation bounds

For contiguous wide inputs in the working dtype, the real cost matrix can be borrowed. Solver state is $O(r + c)$, and pair output is $O(r)$. A tall input may require an $O(rc)$ orientation copy; dtype or layout conversion may also require a dense copy. Including these costs, storage remains $O(rc + r + c)$, rather than requiring a c-square or (r+c)-square cost matrix.

Batch input storage is shared with the call. Each active worker retains its own workspace and can reuse it across problems; results are collected in input order, and the lowest failing index is reported on error. JVRUST and JVHYBRID schedule problems on reusable Rayon thread pools. JVCPP uses a reusable C++ thread pool and claims consecutive small problems in chunks, bounded to about 4096 costs per claim so that adjacent difficult problems are not serialized on one worker. Wide inputs require no per-worker padded cost matrix; tall problems can require orientation buffers. This distinction is relevant to peak memory as worker counts increase.

A conservative operation count also exposes an implementation limitation. Each final augmenting search visits at most r occupied real columns before reaching a terminal, scanning up to c candidates at each step. Across r roots, this gives an $O(r^2c)$ bound for the augmentation portion. If a row-reduction pass permits $O(rB)$ scans with budget parameter B, its cost is $O(rcB)$. The wider-rectangle policy can use $B = c$, whereas specialized policies bound B by a small constant. Thus a conservative bound for the complete rectangular implementation is $O(rcB + r^2c)$, and can be $O(rc^2)$. This is not a claim of an improved worst-case bound over established rectangular assignment methods.

4.5 Numeric and interface contract

All implementations support native float32/float64 storage and explicit precision selection. Required square problems use float64 potentials and distances even for float32 costs; rectangular and finite-limit paths use the selected working precision. Finite conversion overflow is checked.

For a finite limit, the native boundary converts a half-limit and adds it to itself to obtain the working row penalty. In finite precision, $p = 2\,\mathrm{round}_T(L/2)$ need not equal the supplied double-precision L exactly. The theoretical equivalence must therefore be applied to the effective working model. The timed study uses float64 inputs throughout. The independent float32/float64 enumeration study uses binary-rational costs and exactly representable limits.

Positive infinity represents a forbidden edge. NaN and negative infinity in the matrix are outside the documented supported contract. Required matching can be infeasible, and reduced-cost overflow can also cause an error; these conditions should not be conflated. As defensive checks, a path reversal that does not reach its root and a final matching whose row and column mappings disagree are reported as internal errors rather than returned.

All timed methods return int64 pair arrays and omit total-cost calculation. Pair ordering remains native to each interface. Untimed validation checks feasibility and objective values, rather than demanding identical arrays when optimal assignments tie.

5 Experimental design

5.1 Software and environment

The experiments evaluate installed PyPI releases LAPX 0.11.1, lapcpp 1.0.0, and laprust 1.0.0 in two separately prepared environments. Python 3.14.9 uses NumPy 2.5.4 and SciPy 1.18.1; Python 3.11.17 uses NumPy 2.4.6 and SciPy 1.17.1. Runtime records preserve installed package metadata and native-binary hashes, and each suite verifies its binary identities. The measurements were collected on 24 October 2026. The main tables and runtime figures report Python 3.14.9 unless explicitly stated otherwise; Section 6.10 and Appendix A compare the complete Python 3.11.17 replication.

Both environments used the same Intel Core i9-13900KS host, with 16 physical cores and 24 logical processors reported as available to the operating system, on Windows build 26300. These are installed release binaries, not a controlled compiler experiment. Recorded source checkouts support package and native tests; they do not by themselves establish every build flag of a downloaded wheel. The former locally built wheel measurements are not used in this revision.

Single-matrix timings were restricted to logical CPU 0. Batch timings were restricted to logical CPUs 0–3, with one or four requested native workers. Four logical processors are not treated as four physical cores. BLAS/OpenMP pools visible to threadpoolctl were limited to one thread. CPU frequency, turbo behavior, and external activity were not fixed.

The base and extension suites ran sequentially within each environment, in separate collection sessions. Per-logical-CPU utilization samples accompany the runs. Inputs, seeds, adapter semantics, method order, and repetition counts match across environments. NumPy, SciPy, native wheels, and session state also differ, so cross-environment differences cannot be attributed to the Python interpreter alone.

5.2 Equivalent objectives and output work

The seven timed solvers are JVRUST, JVHYBRID, JVCPP, JVX, JVS, JVC, and SciPy. All receive contiguous float64 input. The three implementations of the proposed solver explicitly select float64; JVS disables its default float32 preference.

For required assignment, each public interface receives the real matrix. Its normal rectangular handling is included in timing. For finite limits, JVRUST, JVHYBRID, JVCPP, and JVX use their public limit support. JVS, JVC, and SciPy receive an explicitly augmented (N+M) square matrix, which is constructed inside the timed call. Dummy output pairs are filtered back to the real matrix. Required outputs are normalized to int64 pairs inside timing, and all available total-cost calculations are disabled.

This comparison answers a practical question: what does each configured interface cost when solving the same problem and returning a similar usable result? It is not a measurement of native kernels in isolation. Its finite-limit SciPy configuration uses square expansion. Section 5.6 additionally evaluates two compact formulations, so the square comparison is not treated as the best achievable finite-limit baseline.

5.3 Workloads and timing

Twelve single-matrix cases include uniform square and rectangular inputs, both orientations, a moderately rectangular case, a congested block case, tied costs, and finite positive, tight, negative, and square limits. Each case uses four independently generated matrices. Uniform costs lie in [0,1); the signed case uses [-1,1). Forbidden entries occur with probability 0.5 in the main finite-limit cases and 0.2 in the signed case.

The congested 256×320 case makes three quarters of the rows favor one quarter of the columns. Costs outside that preferred block, and in the opposite block for the remaining rows, receive an increment of 3. The tied case quantizes uniform costs to eighths. These cases test structure beyond independent continuous costs.

Data generation and all reference checks occur outside timing. Each method is warmed on every matrix. A calibration chooses the number of four-matrix cycles for a target 30 ms timing block, with at least one cycle. Nine rounds use a seeded shuffled method order. Garbage collection is disabled only within timed blocks. Results are medians across round averages; interquartile ranges describe timing-round variability rather than the distribution of individual problem instances.

Three batch cases use 64 problems of size 64×80 without a finite limit, 16 problems of size 128×160 with limit 0.5 and 50% forbidden entries, and 1024 problems of size 8×10 without a finite limit. The last case represents many small independent assignments available together. Native JVRUST, JVHYBRID, JVCPP, and JVX batch calls are compared at one and four workers, along with sequential JVRUST and SciPy loops. Batch times include output-list construction and cover the entire batch.

5.4 Correctness validation

An independent memoized enumeration computes optimal objectives for dimensions 0–5 by 0–6, both float32 and float64, ten generated matrices per shape, and five limit settings: infinity, -1, 0, 1, and 2. Costs are integers from -16 through 16 divided by 8, with 20% forbidden entries. Each of the 4,200 problem settings is solved by JVRUST through both laprust and LAPX, by JVHYBRID, and by JVCPP, giving 16,800 calls.

The checks verify pair dimensions, dtype, index bounds, uniqueness, finite selected edges, required cardinality where applicable, and the full penalized objective. Infeasible required problems are identified by enumeration and must produce an error. Binary-rational inputs permit exact objective comparison for this bounded study.

In addition, the installed packages are tested against their own test suites, and the Rust and C++ cores against their native unit tests. Every timed single instance and every batch output is checked against an independently formed SciPy reference objective before measurement.

5.5 Memory measurement

Memory measurements use one fresh process for each method and condition. After imports and a small warmup solve, a contiguous 256×4096 float64 input is created and touched by generation. Windows process counters record the working set and its high-water mark immediately before and after the solve.

The reported additional peak is the increase in the process high-water mark above the input-prepared baseline. It is not an allocation count or total process memory. The original measurements use one trial per condition and are descriptive. Section 5.8 adds five fresh-process trials per condition across a size and orientation sweep, alongside exact transformed-matrix size calculations.

5.6 Implicit rejection against compact formulations

A separate experiment runs JVRUST and JVCPP together with SciPy using (a) the square expansion, (b) r dummy columns, and (c) capped costs as defined in Section 3.5. JVHYBRID shares JVRUST's rectangular and rejection code and is not repeated. All five methods run in the same randomized timing rounds. Conversion, construction of the transformed costs, solving, removal of dummy or capped pairs, orientation mapping where necessary, and int64 pair construction are included. The binaries, CPU affinity, four instances, nine rounds, and 30 ms target blocks are those of the main study; the samples are archived separately, and ratios within this experiment use its own measurements.

The first five workloads reproduce the main study's finite-limit inputs by consuming the same random stream in full. Five additional workloads cover 128×160 inputs at limits 0.005 and 0.001 with 50% forbidden costs, a tall 160×128 case at 0.005 with 50% forbidden costs, a wide 64×512 case at 0.005 with 50%

forbidden costs, and the congested 256×320 block construction at limit 0.5 without forbidden entries. The additional generator uses seed 20261122. These cases were chosen to exercise rejection and structured competition; they are not a held-out random sample of applications.

Before timing, all five methods are compared with independent enumeration on 2,100 float64 problems: dimensions 0–5 by 0–6, ten binary-rational matrices per shape, 20% forbidden entries, and limits -1, 0, 0.125, 1, and 2. This gives 10,500 solver calls, separately from the main validation. Four targeted cases add 20 calls covering all-forbidden costs, equality ties, the Section 3.2 counterexample, and negative costs. Every timed output is checked for feasibility and penalized-objective agreement with the square-expanded reference.

5.7 Matrix-size scaling

A further experiment measures runtime over a predeclared size grid. The presentation follows the dimension sweeps in Figures 4 and 5 of *Linear Sum Assignment with Edition* [5]: both dimensions grow together, or one remains fixed, and the same measurements appear on linear and logarithmic axes. This is a new experiment with the uniform-penalty objective, not a reproduction of that paper's general edit-cost benchmark.

Three shape families use independent uniform [0,1) costs without forbidden entries: square $N = M$ with $N = 16, 32, 64, 128, 256, 512$, and 1024; wide $M = 4N$ with $N = 16$ through 512 in powers of two; and fixed $N = 64$ with $M = 64$ through 4096 in powers of two. Each family is measured for required assignment and for finite-limit assignment at $L = 0.5$. A fourth family applies the congested block construction from Section 5.3 with $M = 5N/4$, $N = 16$ through 1024 in powers of two, and $L = 0.5$. This gives 47 shape/objective settings. The uniform finite-limit sweeps exercise the rejection path but need not produce substantial rejection; the congested sweep separately exercises persistent competition and rejection. It is not claimed to be a worst-case construction.

Required curves compare all seven solvers. Finite-limit curves compare JVRUST, JVCPP, JVX, and the three equivalent SciPy formulations from Section 3.5; JVHYBRID shares JVRUST's finite-limit code. Input and adapter work, int64 outputs, CPU affinity, warmup, four instances, nine shuffled timing rounds, and the 30 ms target block follow the main protocol. At large sizes a single four-matrix cycle can exceed 30 ms; every method still receives nine rounds. Data generation and objective checks are excluded from timing. The input generator uses seed 20261222, and the method-order generator uses 20261223. Each setting receives fresh matrices; a larger input does not contain the smaller input as a submatrix.

The horizontal coordinate is the sum $N + M$ of the original dimensions, not the number of matrix entries or the expanded problem order. Curves connect measured medians; bands and error bars show timing-round interquartile ranges. Both axes are logarithmic in the right panels. No extrapolation or fitted complexity exponent is used. Identical $N + M$ values from different shape families need not represent equal input storage or equivalent problems.

5.8 Larger inputs, special cases, and repeated memory measurements

An additional predeclared grid contains 132 timing cases, 528 method/case settings, and 4,752 timing samples per environment. Large squares have orders 1000, 2000, 3000, and 5000. Large rectangles have shapes 1000×2000, 2000×3000, and 4000×5000, together with their exact transposes. Extreme rectangles pair 64 rows with 4096, 8192, or 16384 columns, again in both orientations. Costs are uniform [0,1); finite-limit large and extreme cases use $L = 2/\max(N, M)$. This size-dependent limit is distinct from the fixed $L = 0.5$ base scaling experiment.

Structured sweeps use 256×320 and 512×640 inputs. Congestion varies the fraction of rows competing for the first quarter of columns over 0.25, 0.5, 0.625, 0.75, and 0.9, with an off-block increment of 3. Forbidden-entry probabilities are 0, 0.5, 0.9, and 0.99; a finite diagonal is planted to guarantee required feasibility. Tied costs use 1, 2, 8, or 256 quantization levels. Their finite limits are 0.5 for congestion, $2/M$ for forbidden entries, and 0.25 for ties. Signed costs use uniform [-1,1), 50% nominal forbidden entries and a finite feasible diagonal on 256×320 and its transpose, with $L = -0.25$ for finite-limit assignment.

Rejection sweeps use 128×160 and 256×320 inputs, either uniform with 50% forbidden entries or congested without forbidden entries. Eleven limits, 0, 0.0005, 0.001, 0.002, 0.005, 0.01, 0.02, 0.05, 0.1, 0.5, and 4, reuse the same costs within a family. Changing the limit changes the objective, so this experiment does not isolate a causal effect of the rejected fraction. The seed is 20261024, with stable family and instance keys. Runtime transposes use identical costs in transposed orientation. Each case retains four inputs, nine shuffled rounds, and the 30 ms calibration target. Every timed output is validated outside timing, and hashes are checked again to detect input mutation.

The focused required comparison uses JVRUST, JVHYBRID, JVCPP, and native rectangular SciPy; the finite-limit comparison uses JVRUST, JVCPP, capped SciPy, and dummy-column SciPy. JVX, JVS, and JVC are excluded from this extension to keep the larger experiments tractable; the original all-solver comparisons remain in Figures 1–6. Thus a winner in a focused plot means a winner among those selected methods. The declared resource limits are 240 seconds per method/case, 1200 seconds per complete case, and 120 seconds per memory trial. No method was excluded and no measurement failed in either completed extension.

Repeated memory trials use $256 \times \mathrm{M}$ for $M = 256$, 512, 1024, 2048, and 4096, plus the nonsquare transposes. Required methods are JVRUST, JVCPP, SciPy, and JVX; finite-limit methods are JVRUST, JVCPP, capped, dummy-column, and square-expanded SciPy at $L = 0.5$. There are 81 conditions and five fresh processes per condition, totaling 405 trials per environment. A small warmup precedes input creation; validation occurs after peak capture. Each shape/trial is paired across methods, while memory orientations use independently generated, distribution-matched arrays to avoid a pre-baseline transpose allocation. Medians and interquartile ranges include process and input variation and are not allocation counts.

Finally, 23 edge fixtures generate 129 solver checks, including 15 expected infeasibility checks. They cover empty dimensions, one row or column, all-forbidden and all-zero costs, costs equal to the rejection limit, signed costs, an isolated infeasible row, and neighboring floating-point values around the limit. The small fixtures use independent exhaustive objectives; tied pair identities need not agree.

5.9 Cross-environment comparison

Each environment contributes 994 method/case settings across 204 cases: 12 single-call, three batch, ten formulation, 47 scaling, and 132 extended cases. All 8,946 raw timing samples per environment are retained. The replication checks recorded input hashes and schedules, then independently recalculates medians and quartiles from the samples. It also compares per-input objectives, cardinalities, and repeated-memory summaries.

Ratios in the manuscript divide within-run median times; they do not average ratios from different sessions. Timing rounds are not paired across environments. Some supplementary focused charts summarize within-round paired ratios and are explicitly labeled that way. Interquartile ranges describe the recorded timing rounds or memory trials, not population confidence intervals. The workload grid is designed rather than sampled from applications; no population significance tests or universal interpreter-speed claim are made.

6 Results

6.1 Correctness outcomes

All four solver entry points agree exactly with enumeration on all 16,800 calls: JVRUST through laprust and through LAPX, JVHYBRID, and JVCPP. Eighty calls correspond to infeasible required problems and correctly produce errors. The maximum absolute objective discrepancy on feasible settings is zero. This result applies to the finite grid and shapes described above.

The installed packages pass their Python test suites: 2,704 LAPX, 1,167 lapcpp, and 1,167 laprust tests. The corresponding source checkouts pass 24 Rust unit tests for laprust and 21 for LAPX's solver core with all features enabled, and lapcpp's native C++ test program. Both environments pass these same suites. These

suites include numeric contract cases, primal-dual certificates, coexistence and concurrency checks, fault injection, and comparisons between completion implementations. They support implementation confidence but do not constitute a proof for arbitrary floating-point inputs or every platform.

All 336 single-instance benchmark checks and 11,040 batch-output checks pass in each environment. The extension adds 2,112 timed-input objective checks and 129 edge checks per environment, all passing. Across the formulation, scaling, and extended suites, all 3,520 recorded objectives and match counts agree exactly between environments. Pair-array equality is not a required outcome in tied problems.

6.2 Single-matrix runtime

Table 3 reports median microseconds per matrix. Figure 1 shows the same cases with timing-round interquartile ranges. JVS, JVC, and SciPy values in finite-limit rows include explicit augmentation.

Table 3: Median runtime in microseconds, nine rounds and four matrices per case. Smaller is better. Raw samples and full interquartile ranges accompany the manuscript.

Case	JVRUST	JVHYBRID	JVCPP	JVX	JVS	JVC	SciPy
Square 128×128	68.47	69.82	78.82	127.52	120.93	296.40	294.05
Square 512×512	1108.85	1123.00	1190.53	3385.24	3504.47	6294.77	5247.77
Wide 64×512	13.01	12.56	15.02	957.28	890.37	1954.62	76.29
Tall 512×64	26.01	26.20	86.26	953.45	874.88	1872.11	158.69
Rectangle 256×320	76.54	77.29	87.25	188.29	187.12	4379.26	588.21
Blocks 256×320	9707.80	9776.38	10714.25	10442.83	10583.38	6280.60	1565.43
Ties 128×160	56.20	61.41	59.67	68.46	65.97	104.75	84.46
128×160, $L = 0.5$	14.89	15.39	21.17	551.83	483.97	2292.61	447.76
256×320, $L = 0.5$	76.52	77.20	84.92	4053.53	3381.00	15973.10	2315.12
Tight 128×160, $L = 0.05$	9.96	10.43	15.53	487.20	417.84	2656.32	1169.17
Signed 128×160, $L = -0.25$	14.83	16.03	21.45	1095.88	957.44	3292.79	448.89
Square 128, $L = 0.5$	91.49	95.30	107.96	213.91	207.66	656.27	563.53

An implementation of the proposed solver is fastest among the configured methods in six of the seven required cases and in all five finite-limit cases. On the uniform 256×320 problem, JVRUST takes 76.54 microseconds versus 588.21 for SciPy. The tall case costs more than the corresponding wide case for JVRUST and SciPy, consistent with orientation-dependent work, although this experiment does not isolate every source of the difference.

The structured cases qualify the overall result. On congested blocks, JVRUST requires 9.71 ms while SciPy requires 1.57 ms: SciPy is about 6.2 times faster, and JVC is also faster than all three implementations. On tied costs, JVRUST and JVS are close: JVRUST takes 56.20 microseconds and JVS 65.97. Matrix dimensions alone do not predict the fastest method. Initialization and reduction behavior on the cost structure also matter.

Finite-limit JVRUST runtimes range from 9.96 to 91.49 microseconds. The 128×160, $L = 0.5$ case is 30.1 times faster than the square-augmented SciPy configuration. This ratio applies to that interface/formulation comparison; Section 6.5 reports a separate paired comparison against smaller and unexpanded alternatives.

Most finite-limit timing instances still have a complete smaller-side matching at the optimum. The tight $L = 0.05$ case averages 119.75 pairs, leaving 8.25 of 128 smaller-side rows unmatched. Thus the principal measured benefit in the other finite-limit cases is avoiding the augmented formulation, rather than avoiding a large fraction of real matches. The added rejection and structure sweeps broaden the timed evidence (Section 6.8), while exhaustive validation remains restricted to small inputs.

The three implementations differ in ways that follow their structure. JVRUST and JVHYBRID are within a few percent on the required square cases, so the C++ completion kernel gives no consistent advantage; on their shared Rust paths, small differences are not evidence of different algorithmic capabilities. JVCPP is 6–15% slower than JVRUST on the required cases other than the tall input, 1.1–1.6 times slower on the

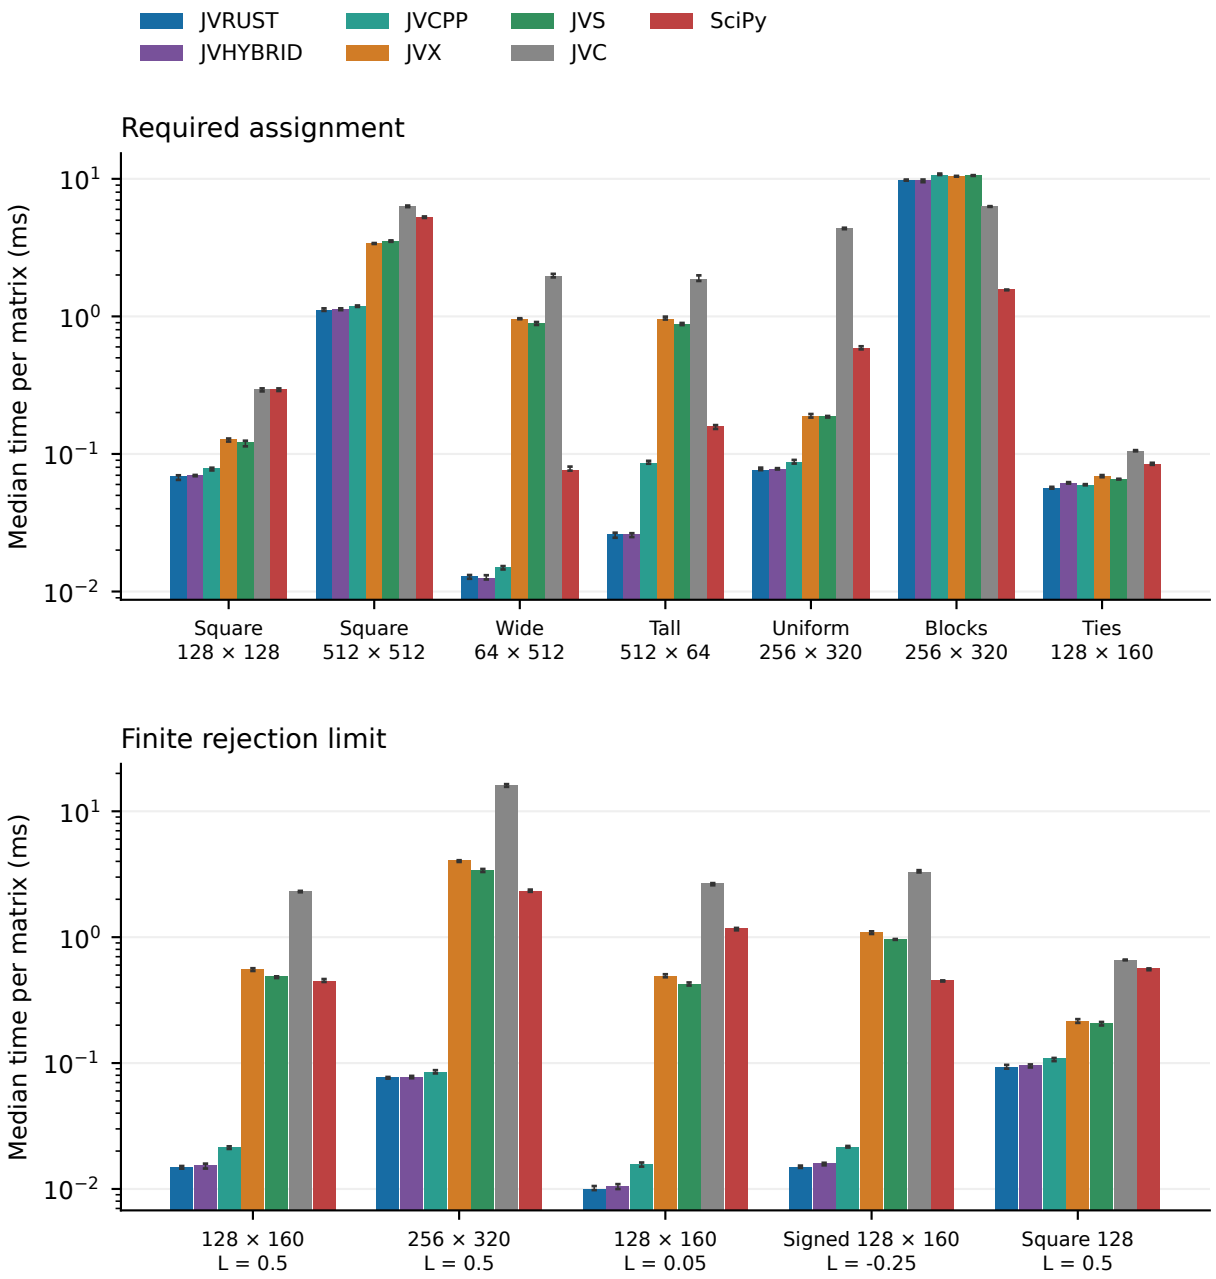


Figure 1: Required and finite-limit assignment runtimes on a logarithmic scale. Error bars show the 25th–75th percentiles across timing rounds. The block and tied cases are retained rather than reporting only favorable workloads.

finite-limit cases, and 3.3 times slower on the tall 512×64 case. The larger tall-input gap is consistent with orientation or other implementation overhead, but these timings do not isolate its cause. The tall and wide cases use independently generated matrices; the paired transpose experiments in Section 6.7 broaden the orientation evidence, but phase-level profiling is still needed to separate copying, search, and output costs.

6.3 Batch runtime

Table 4 compares whole-batch latencies. Four-worker native calls improve throughput within the restricted four-logical-processor affinity, but the amount depends on the workload.

JVRUST's four-worker call is 2.32 times faster than its one-worker call on the 64×80 batch, 1.93 times faster on the finite-limit batch, and 2.13 times faster on the 1024-problem batch. With larger problems, serial batching is close to the JVRUST loop, and the benefit comes from parallel execution with retained native

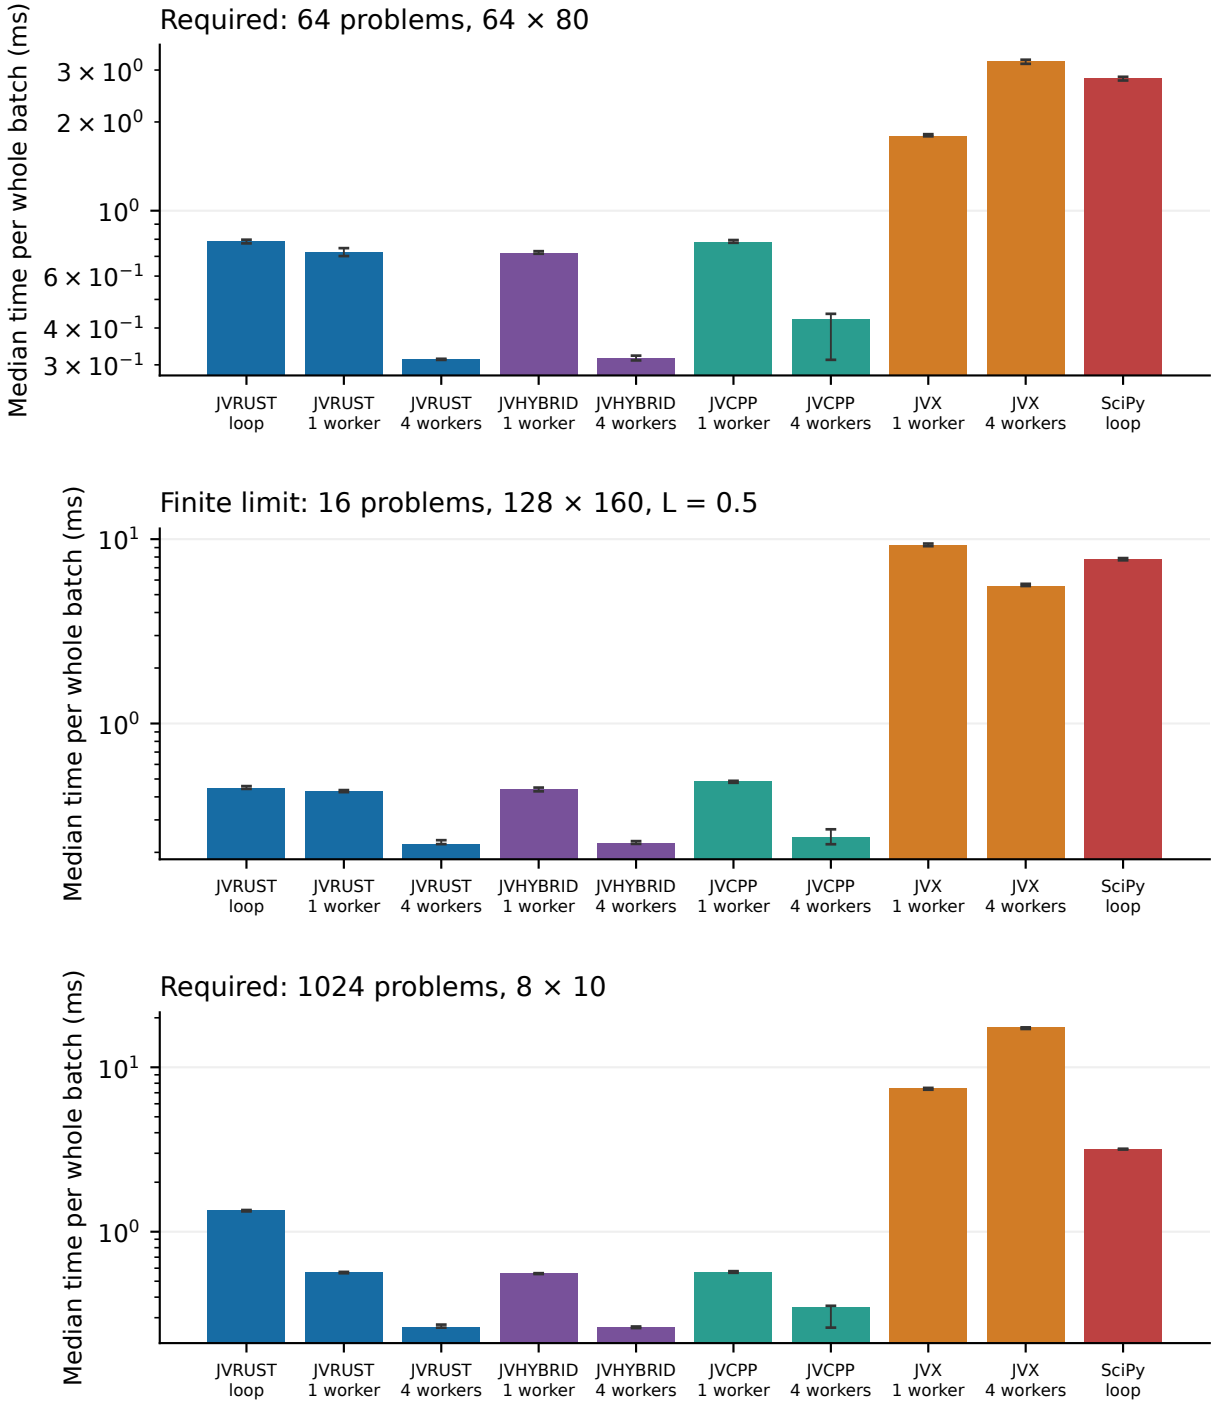


Figure 2: Whole-batch runtimes and timing-round interquartile ranges for one and four requested workers. All batch methods run within logical CPUs 0–3.

Table 4: Median milliseconds per complete batch. Finite-limit SciPy calls include augmentation. All methods return usable pair arrays, and reported times exclude input generation.

Method	**$64 \times (64 \times 80)$ required, ms**	**$16 \times (128 \times 160)$ $L = 0.5$, ms**	**$1024 \times (8 \times 10)$ required, ms**
JVRUST sequential loop	0.789	0.444	1.343
JVRUST batch, 1 worker	0.725	0.429	0.563
JVRUST batch, 4 workers	0.313	0.223	0.264
JVHYBRID batch, 1 worker	0.718	0.442	0.556
JVHYBRID batch, 4 workers	0.316	0.225	0.263
JVCPP batch, 1 worker	0.779	0.486	0.565
JVCPP batch, 4 workers	0.430	0.242	0.346
JVX batch, 1 worker	1.795	9.220	7.413
JVX batch, 4 workers	3.213	5.580	17.429
SciPy sequential loop	2.824	7.823	3.181

workspaces. With 1024 problems of size 8×10, the batch call is 2.4 times faster than the loop even with one worker, consistent with amortizing Python-call and output-handling overhead; four workers make it 5.1 times faster than the loop. JVHYBRID follows JVRUST closely. JVCPP's one-worker batch calls are slower than JVRUST's in every case. With four workers, the 64×80 batch medians differ appreciably in Python 3.14: JVCPP takes 0.430 ms versus JVRUST's 0.313 ms, approximately 37% longer. Its timing-round interquartile range is wide, 0.312–0.447 ms. The Python 3.11 replication does not preserve this gap (Section 6.10). JVCPP is slower on the other two batches, most on the 1024-problem batch, where its four-worker call takes 0.346 ms against 0.264 ms.

The JVX four-worker call is slower than its serial call on both required batches, while improving the finite-limit batch. The result demonstrates that requesting more workers is not inherently beneficial. This study does not isolate scheduling, Python dispatch, and memory effects sufficiently to assign that regression to one cause.

6.4 Peak working set

Table 5 reports additional peak working set above the input-prepared baseline. The common 8 MiB input is already resident. Figure 3 plots these increments.

Table 5: High-water-mark increments for 256×4096 float64 input, one fresh process per condition. These values are not total process memory.

Method	**Required, MiB**	**Finite $L = 0.5$, MiB**
JVRUST	0.125	0.168
JVHYBRID	0.117	0.152
JVCPP	0.141	0.133
JVX	128.145	144.660
SciPy	0.215	144.746

The measured increments agree closely with the 128 MiB padded and 144.5 MiB augmented cost-matrix calculations, plus smaller workspace effects. The three implementations add 0.12–0.17 MiB in both conditions. Required SciPy is similarly small, confirming that this memory benefit is shared by existing native rectangular methods.

The small increments do not imply zero allocation. They are process-level measurements affected by prior high-water marks, allocator reuse, page residency, and the single-trial design. Their main interpretation is the absence of the much larger transformed dense cost buffer in this wide-input experiment.

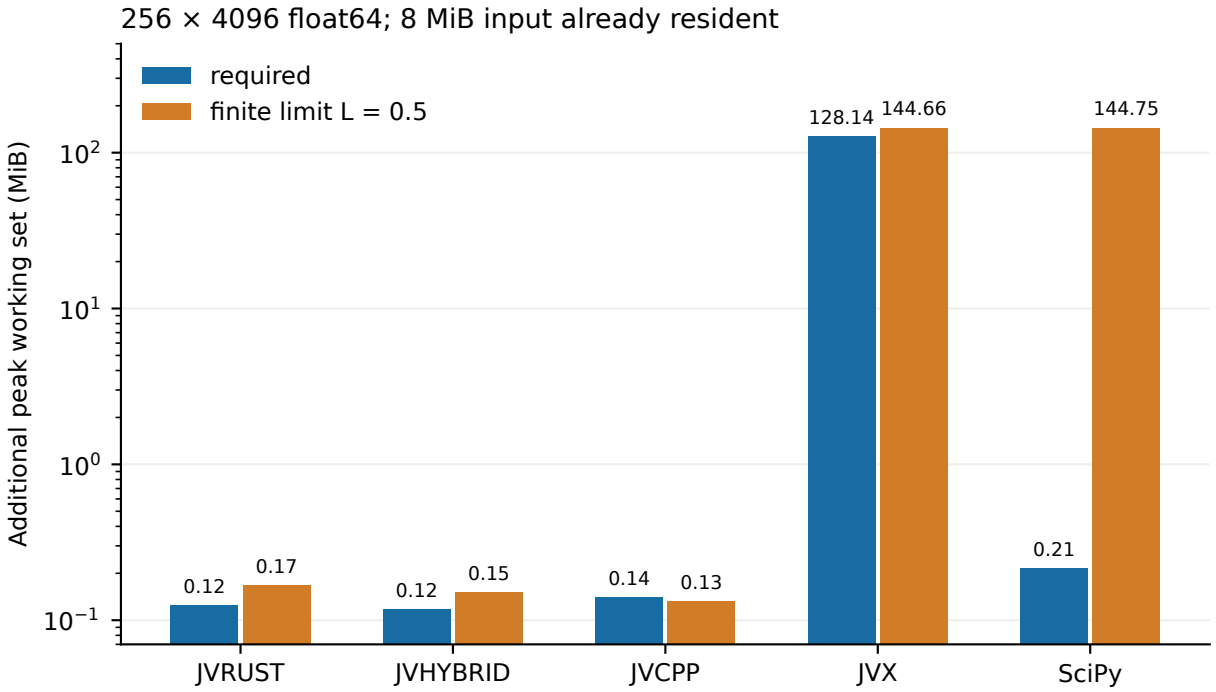


Figure 3: Additional peak working set on a logarithmic scale. The three implementations and required SciPy operate without a transformed dense cost matrix. Finite-limit SciPy uses the explicit augmented baseline defined in Section 5.

6.5 Implicit rejection against compact formulations

The additional enumeration has zero objective error on all 10,500 grid calls, and the 20 targeted checks pass. All 200 timed-instance checks also pass. Table 6 reports the paired run. The mean unmatched count refers to smaller-side rows under JVRUST; equality ties may permit alternative optimal cardinalities.

Table 6: Finite-limit formulation comparison, median microseconds per matrix. Square, Dummy cols, and Capped all use SciPy with the corresponding cost construction. Unmatched is a mean count over four matrices, not a time. Nine timing-round samples and their interquartile ranges are archived for every method/case.

Case	Unmatched	JVRUST	JVCPP	Square	Dummy cols	Capped
128×160, $L = 0.5$	0.00	15.09	22.29	449.53	270.09	136.02
256×320, $L = 0.5$	0.00	78.07	85.79	2377.34	1417.70	606.58
128×160, $L = 0.05$	8.25	10.58	16.07	1206.82	275.33	117.73
Signed 128×160, $L = -0.25$	0.00	15.08	20.95	442.48	241.04	142.01
Square 128, $L = 0.5$	0.00	89.22	104.05	519.32	411.31	223.18
128×160, $L = 0.005$	87.75	6.71	7.80	928.85	186.24	81.58
128×160, $L = 0.001$	117.00	6.16	7.00	524.88	169.16	70.65
Tall 160×128, $L = 0.005$	89.50	14.22	14.89	923.88	177.90	82.48
Wide 64×512, $L = 0.005$	18.00	10.04	10.26	2258.12	219.39	126.05
Blocks 256×320, $L = 0.5$	112.00	1802.40	1884.59	2666.16	1699.89	749.04

For the 128×160, $L = 0.5$ workload, JVRUST is 9.0 times faster than capped SciPy in this run, compared with 29.8 times faster than square-expanded SciPy. The smaller dummy-column construction falls between the two SciPy alternatives. Thus a substantial part of the square comparison is formulation overhead, while a runtime difference remains against an unexpanded baseline on this workload.

JVRUST is faster than capped SciPy in nine of ten finite-limit cases, by factors of 2.5–12.6 on those cases; JVCPP is also faster than capped SciPy in the same nine cases. The congested block case reverses the result: capped SciPy takes 0.749 ms and JVRUST takes 1.802 ms, making capped SciPy 2.4 times faster. This reversal occurs despite both methods avoiding dimensional expansion. The rejection-heavy cases leave

87.75 and 117 of 128 rows unmatched on average, extending the timed evidence beyond nearly complete matchings. These are descriptive ratios of within-run medians, not confidence intervals or an estimate of general application speedup.

The alternatives also change the storage comparison. For a 256×4096 float64 problem, square expansion requires a 144.5 MiB transformed cost matrix, the r-dummy-column construction requires 8.5 MiB, and capped costs require 8 MiB. These are exact buffer-size calculations, not new process-memory measurements. For a suitable wide input, the proposed solver borrows the original costs and needs $O(r + c)$ solver state; capped SciPy avoids the square expansion but our adapter creates an additional original-size cost array. The distinction is avoiding a transformed copy, as well as the number and organization of columns scanned.

6.6 Matrix-size and aspect-ratio scaling

Figures 4 and 5 show the uniform required and finite-limit sweeps, respectively. Figure 6 shows the congested finite-limit sweep with the same plotting convention. Every curve comes from the separate paired scaling run; ratios are calculated within that run. All 1,208 timed-instance feasibility and objective checks pass, and 2,718 timing-round samples are archived across 302 method/setting combinations.

The shape sweeps distinguish increasing the smaller dimension from increasing only the larger one. At 1024×1024 without a finite limit, JVRUST takes 4.410 ms, JVHYBRID takes 4.399 ms, JVCPP takes 4.684 ms, and SciPy takes 24.703 ms. At 64×4096, JVRUST takes 0.070 ms and SciPy takes 0.555 ms, while JVX takes 71.956 ms. Thus a large sum of dimensions does not imply a difficult problem for every implementation. Padding changes both the represented dimensions and the search work; the present measurements do not isolate those effects individually.

For finite-limit 512×2048 inputs at $L = 0.5$, JVRUST takes 0.369 ms, capped SciPy takes 4.281 ms, square-expanded SciPy takes 28.212 ms, and JVX takes 1392.663 ms. The large difference between the two square-expanded implementations shows that transformed dimensions alone do not predict runtime. In these uniform sweeps, which contain no forbidden entries, JVCPP takes up to 1.6 times as long as JVRUST; at the largest size of each family, their median runtimes differ by at most 20% relative to JVRUST, and neither implementation is faster at every point. The narrowing gaps are consistent with a contribution from fixed setup costs, but the measurements do not separate setup from search time. All uniform finite-limit scaling instances have complete smaller-side matchings under JVRUST, so those curves primarily characterize finite-limit handling on easy-to-match random costs. The congested sweep tests a different regime.

The congested curve changes the ordering as size grows. JVRUST is faster than capped SciPy at $N = 16$, 32, and 64; capped SciPy is faster at $N = 128$, 256, 512, and 1024. At 1024×1280, JVRUST takes 167.616 ms and capped SciPy takes 17.789 ms, a 9.4-fold advantage for capped SciPy. JVRUST leaves 43.75% of the smaller side unmatched at every tested congested size, including 448 of 1024 rows at the largest point. Substantial rejection therefore does not itself imply an advantage for the native terminal strategy. Both methods avoid dimensional expansion, and the observed crossover motivates investigation of their initialization and search policies.

6.7 Larger matrices and extreme aspect ratios

The focused comparisons preserve the native-versus-SciPy advantage on the tested large uniform matrices in both environments. Table 7 includes representative large and signed cases from both runs. At 5000×5000 without rejection, JVRUST takes 257.35 ms in Python 3.14 and 253.86 ms in Python 3.11; SciPy takes 1083.77 and 1095.19 ms. At the size-dependent finite limit $L = 0.0004$, the same-sized costs define a different optimization problem: JVRUST takes 11.42 and 11.41 ms, compared with 146.47 and 132.90 ms for capped SciPy. These required and finite-limit times are not an ablation of rejection overhead.

On the large square and large rectangle grids, JVCPP/JVRUST median-time ratios span 0.961–1.061 in Python 3.14 and 0.962–1.058 in Python 3.11. The implementations are therefore within about 6.1% on this grid, with neither winning every point. This observation is specific to the added large cases: the original

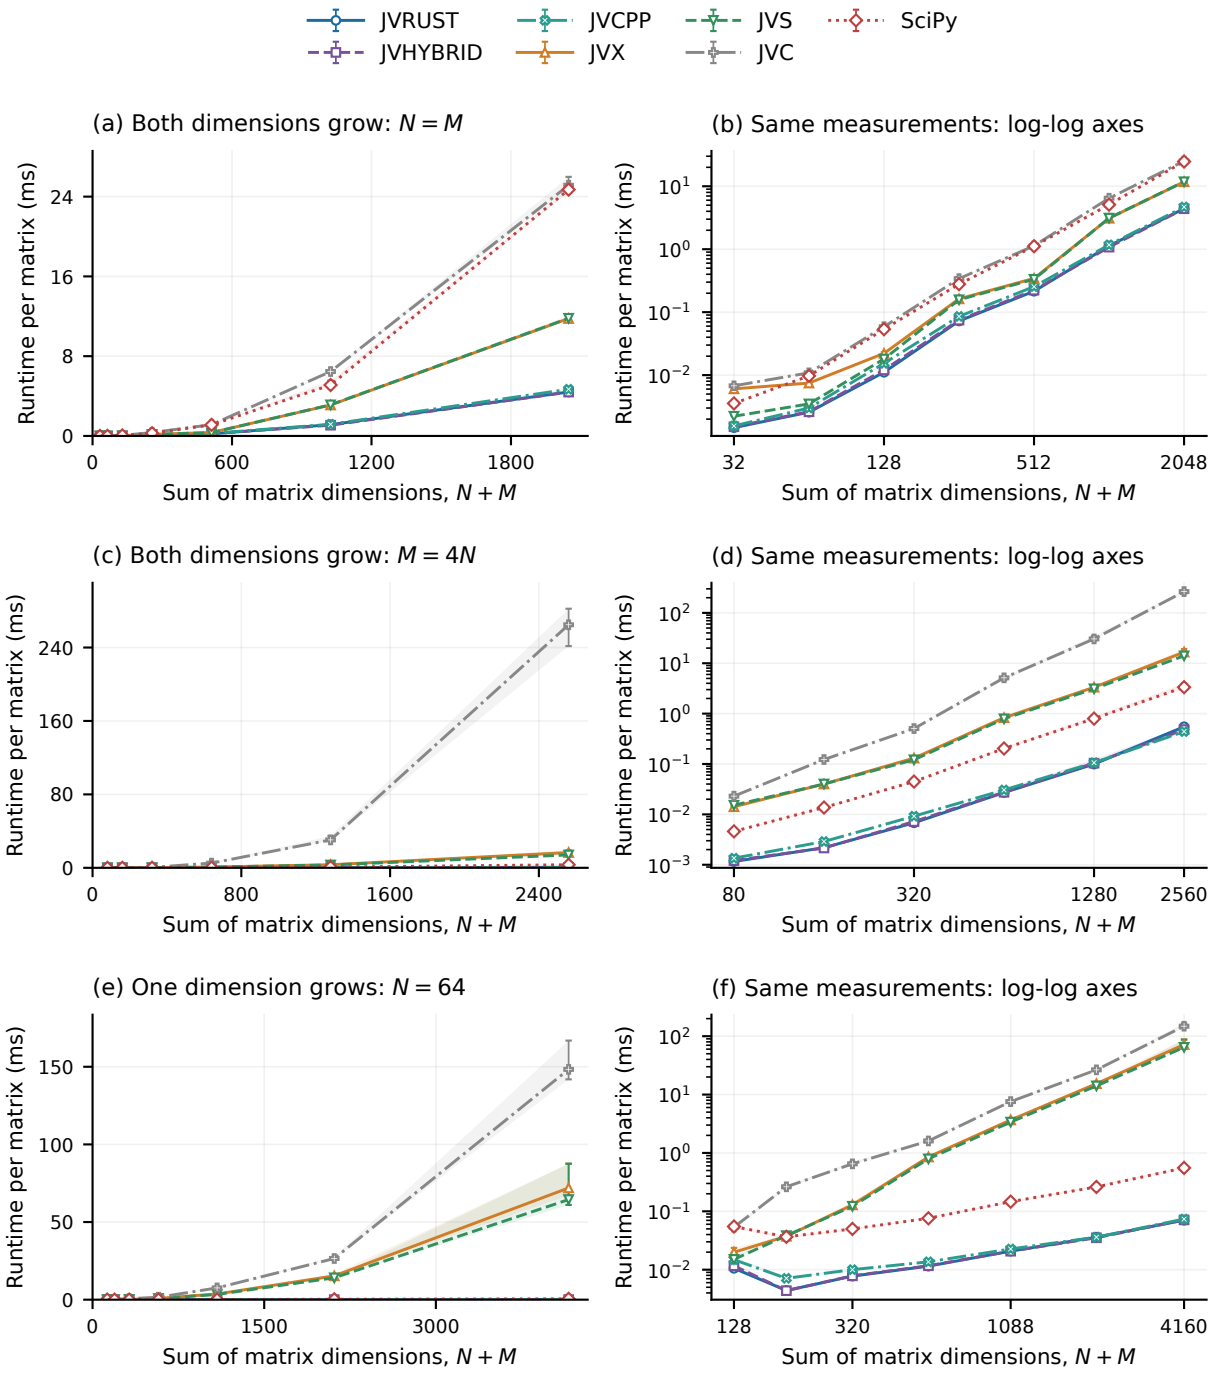


Figure 4: Required-assignment runtime as a function of the sum of the original matrix dimensions. Rows show square matrices, 4:1 wide matrices, and fixed $N = 64$ with increasing M. Left panels use linear axes; right panels repeat the same measurements on log-log axes. Each point is the median of nine timing-round averages over four matrices. Bands and error bars show interquartile ranges. Costs are independent uniform [0,1).

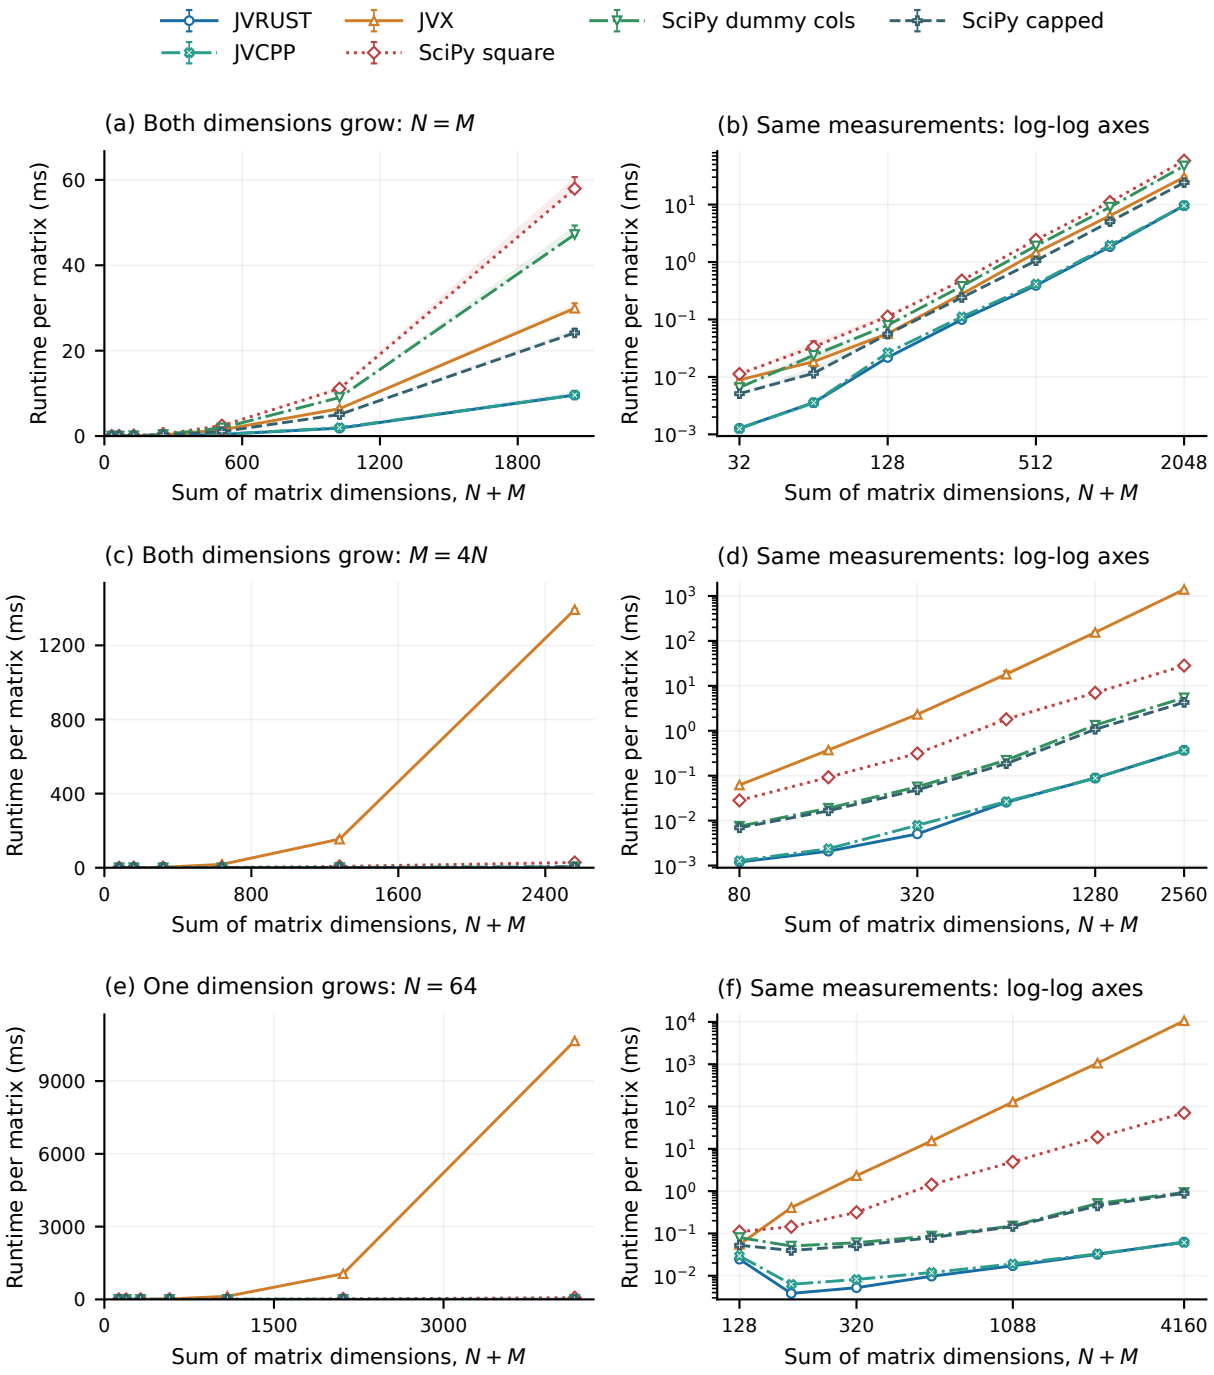


Figure 5: Finite-limit scaling for the same three shape families at $L = 0.5$. Square, dummy-column, and capped-cost curves use SciPy with equivalent objectives and include construction and output filtering. JVHYBRID shares JVRUST's finite-limit code and is omitted. Linear and log-log views repeat the same data; bands and error bars show timing-round interquartile ranges. These uniform inputs are distinct from those in Figure 4.

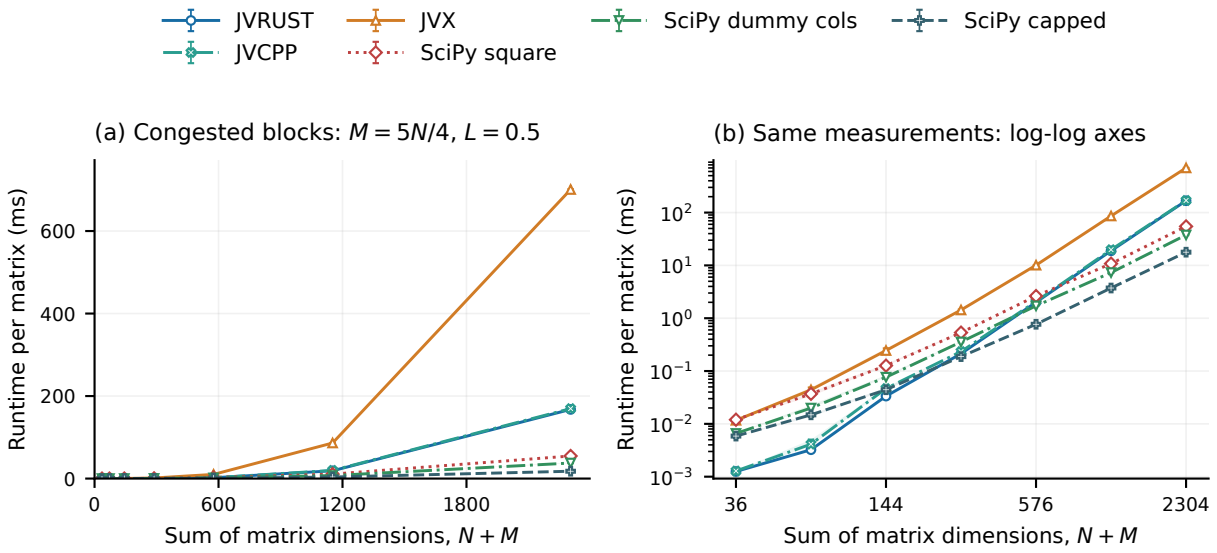


Figure 6: Finite-limit scaling under congested costs, with $M = 5N/4$ and $L = 0.5$. Three quarters of the rows compete for one quarter of the columns; off-block costs receive an increment of 3. Linear and log-log views show the same measured medians and timing-round interquartile ranges. This structured family tests whether behavior seen with independent uniform costs persists under competition.

Table 7: Representative extended-case medians in milliseconds. R denotes required assignment; F denotes the finite-limit regime. The reference is native SciPy for R and capped SciPy for F. Large F cases use $L = 0.0004$; signed F cases use $L = -0.25$. Each value summarizes nine timing rounds over four inputs in its own environment.

Case	JVRUST 3.14	Ref. 3.14	JVRUST 3.11	Ref. 3.11
5000 × 5000 R	257.352	1083.767	253.856	1095.185
5000 × 5000 F	11.424	146.471	11.409	132.900
4000 × 5000 R	14.068	131.876	13.805	137.624
4000 × 5000 F	8.503	102.012	8.489	96.458
Signed 256 × 320 R	0.084	0.810	0.076	0.835
Signed 256 × 320 F	0.063	0.469	0.063	0.496
Signed 320 × 256 R	0.120	0.977	0.109	0.933
Signed 320 × 256 F	0.096	0.575	0.096	0.598

uniform sweep has a larger maximum gap at the largest point of its shape families. The hybrid completion kernel again gives no uniform benefit.

Extreme tall inputs retain a different implementation pattern: JVCPP/JVRUST ratios range from 3.00 to 3.84 in Python 3.14 and 2.98 to 3.81 in Python 3.11. The signed cases in Table 7 also show larger tall than wide times for JVRUST in both regimes. Paired runtime transposes establish that the orientation differences persist on identical cost values, but do not identify a single cause in orientation conversion, search, or output work.

Dummy-column SciPy can be faster than capped SciPy on extreme tall finite-limit inputs. At 16384×64 it takes 5.368 versus 8.141 ms in Python 3.14 and 5.632 versus 8.216 ms in Python 3.11. JVRUST remains faster than both.

6.8 Structure, rejection, and edge cases

The structured sweeps reproduce the central congestion counterexample. At both tested shapes, competing-row fraction 0.25 favors the native implementations over the matching SciPy baseline; fractions 0.5, 0.625, 0.75, and 0.9 favor SciPy in both required and finite-limit regimes. This accounts for sixteen of the extension's

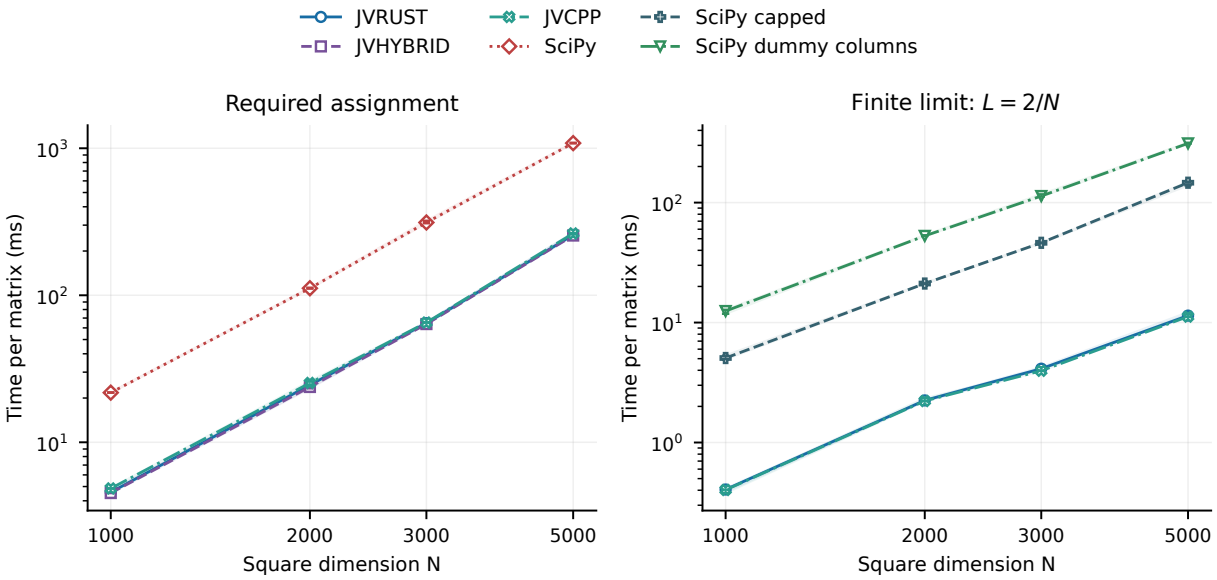


Figure 7: Large uniform squares in Python 3.14. Required and finite-limit panels use different objectives; finite limits are $2/\max(N, M)$. Lines and error bands show medians and timing-round interquartile ranges for the fixed focused method sets.

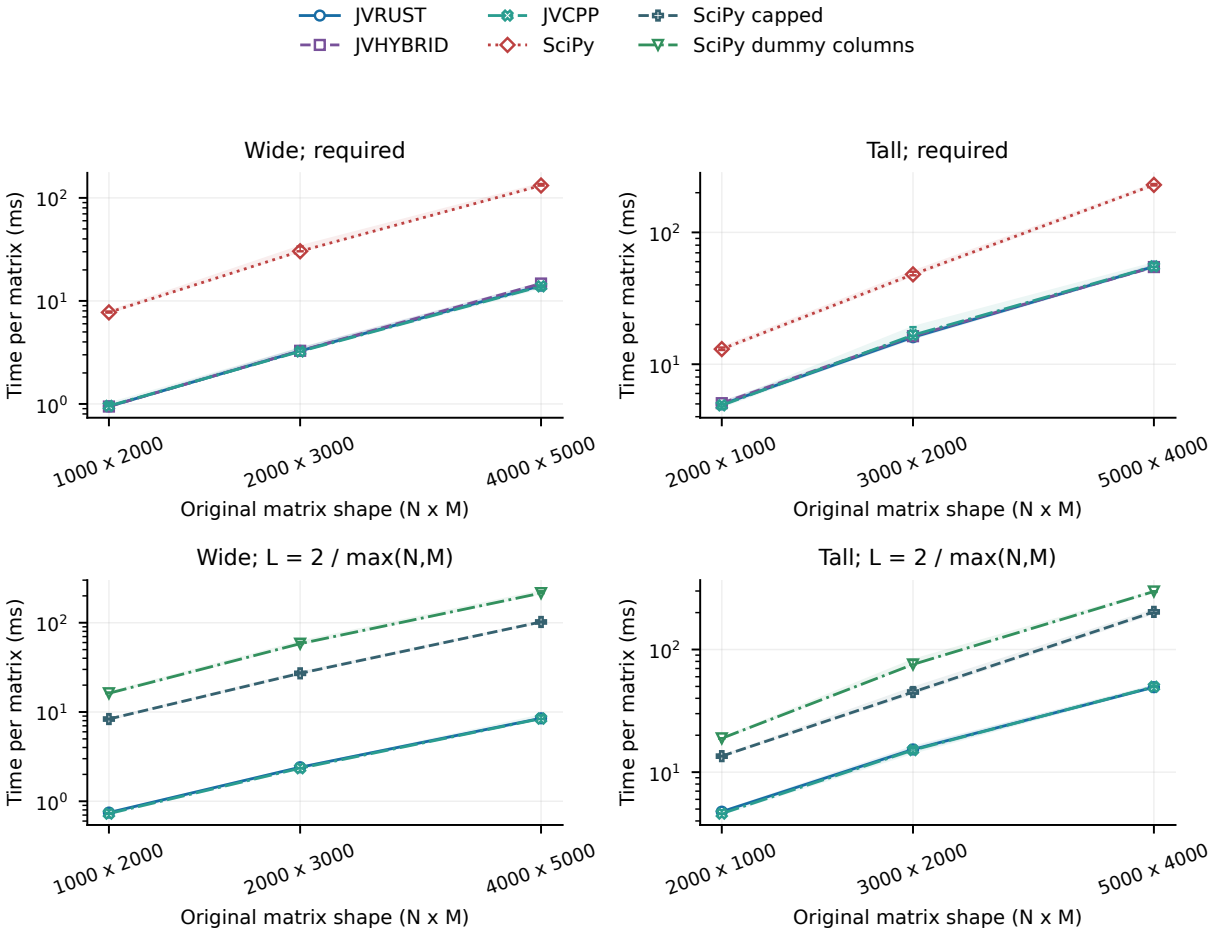


Figure 8: Large rectangles and their exact transposes in Python 3.14, including 4000×5000 and 5000×4000. Required and finite-limit comparisons retain the compact SciPy alternatives. The Python 3.11 replication preserves each native method's ordering against the matching SciPy baseline.

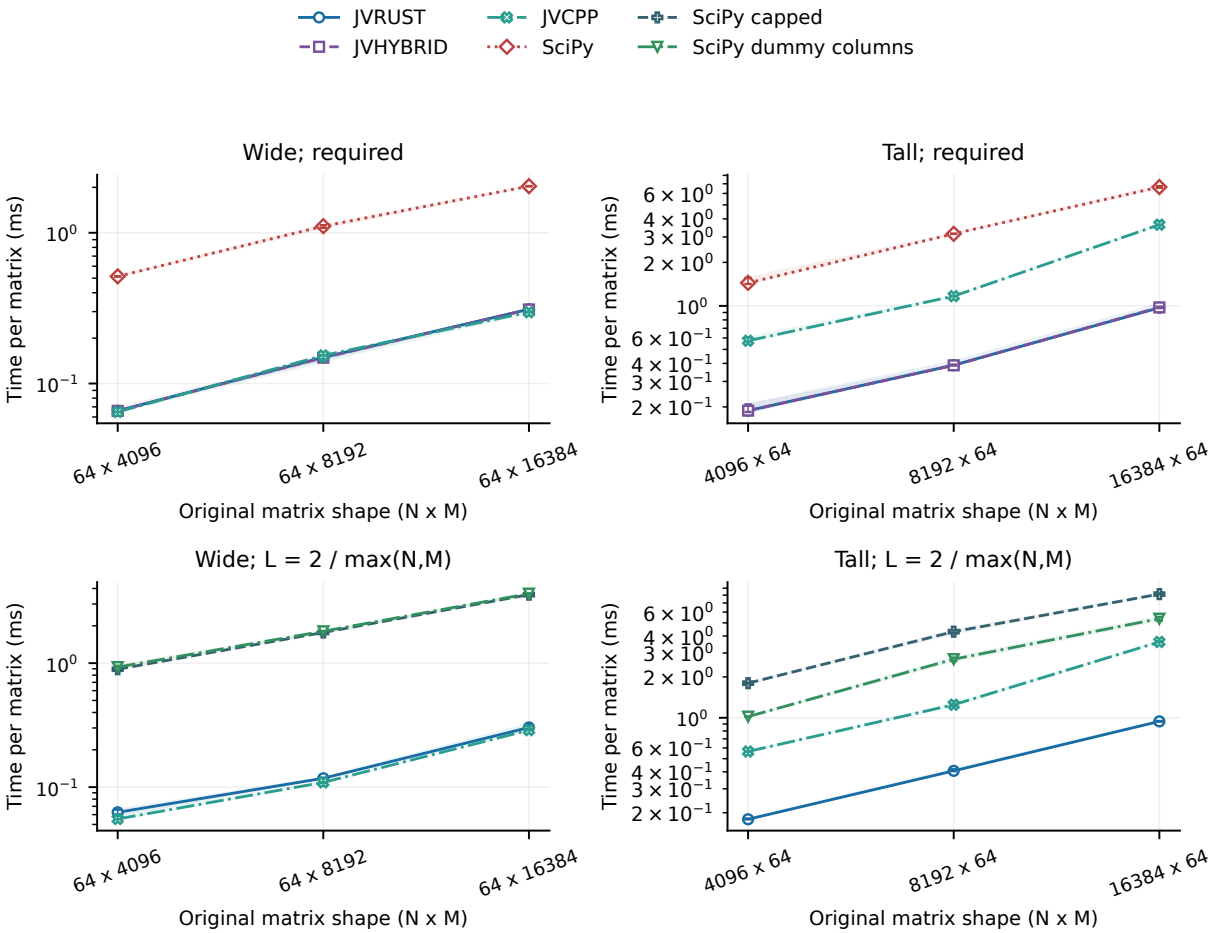


Figure 9: Extreme shapes with smaller dimension 64 and larger dimension up to 16384, in both orientations, in Python 3.14. Transposes reuse the same cost values; finite-limit panels use $L = 2/\max(N, M)$. Interquartile ranges describe the nine timing rounds.

twenty native-versus-SciPy losses in each environment. The sampled fractions do not locate a continuous crossover threshold.

The native methods retain their advantage over the matching SciPy baseline throughout the tested forbidden-entry, tie, and signed grids in both environments. This includes dense representations with 99% nominal forbidden entries and completely tied inputs; it does not compare sparse data structures or establish performance for arbitrary numerical ranges. Distinct selected pairs can still be optimal under ties.

The congested sweeps have roughly 43.75% unmatched smaller-side rows over intermediate limits and full matching at $L = 4$. Capped SciPy leads at the tested limits 0.5 and 4 for both shapes, accounting for the other four native losses; substantial rejection alone does not guarantee a runtime advantage. The limit changes both the optimal matching and the search problem. The separate edge fixtures pass in both environments, including the expected infeasible cases and floating-point neighbors of the penalty boundary; this bounded validation is not a general numerical stability proof.

6.9 Repeated memory measurements

The five-trial sweep confirms both formulation and orientation effects. At 256×4096 and $L = 0.5$, median additional peak working set in Python 3.14 is 0.152 MiB for JVRUST, 8.098 MiB for capped SciPy, 8.641 MiB for dummy-column SciPy, and 144.734 MiB for square-expanded SciPy. The corresponding Python 3.11 medians are 0.137, 8.090, 8.664, and 144.742 MiB. The maximum absolute difference between the two environments' medians over all 81 conditions is 0.08984 MiB.

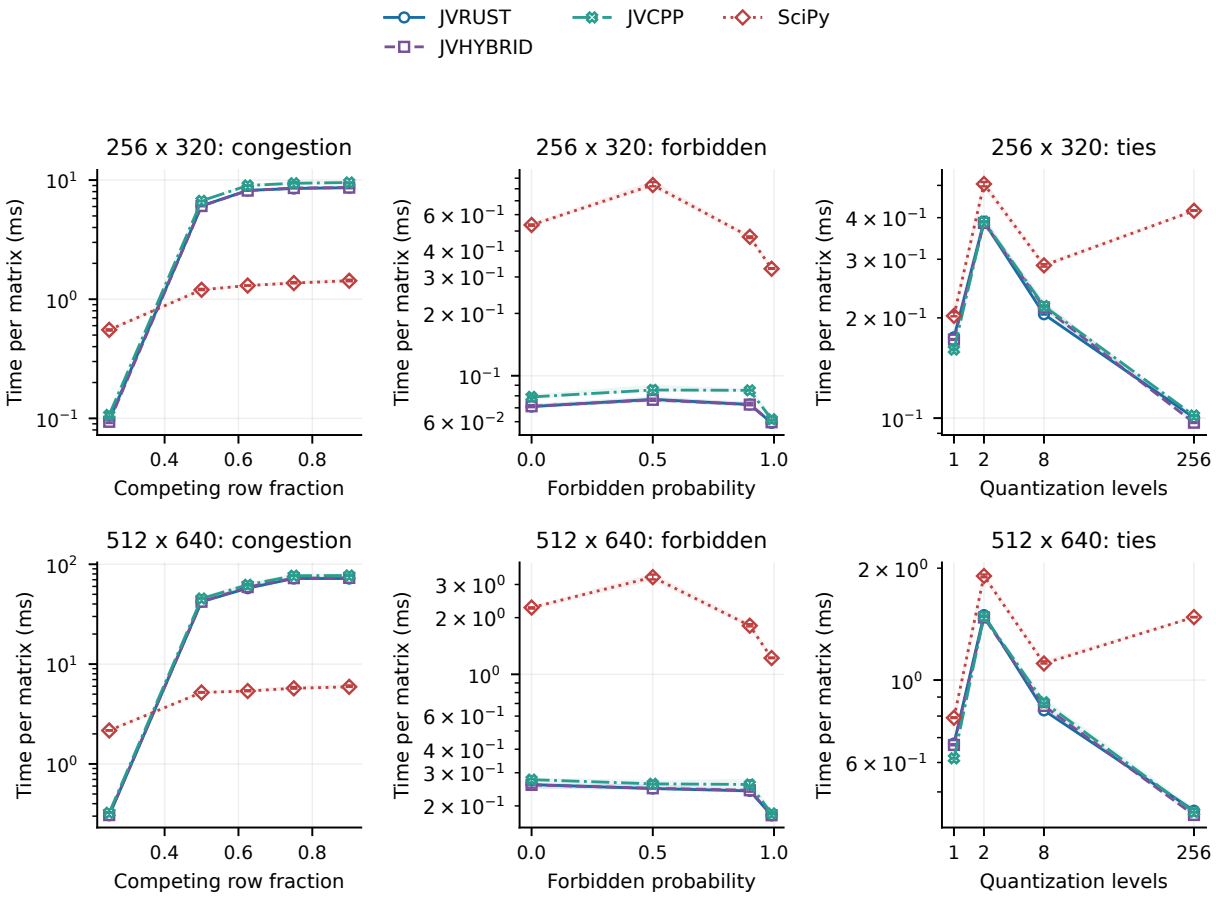


Figure 10: Required structure sweeps in Python 3.14 at 256 × 320 and 512 × 640. The columns vary congestion, forbidden-entry probability, and quantization. A finite diagonal preserves feasibility in the forbidden sweep. All panels use the same four required-assignment methods.

The small wide-input native footprint must not be generalized to tall inputs. For 4096 × 256 at $L = 0.5$, JVRUST's medians are 8.172 and 8.141 MiB in Python 3.14 and 3.11; capped SciPy's are 16.082 and 16.160 MiB, while dummy-column SciPy's are 8.645 and 8.664 MiB. Our capped adapter first creates a capped array in the original tall orientation, after which SciPy allocates a transposed array of the same size; the dummy-column adapter orients before constructing its buffer. This additional original-size allocation is specific to the evaluated adapter path; capped-cost assignment can instead use a single oriented buffer. Square-expanded SciPy remains near 144.7 MiB. These observations are consistent with additional orientation-related work and storage, but the process counters do not trace allocations. Near-zero increments should be compared in absolute units, because percentage ratios can become unstable or undefined.

6.10 Consistency across Python environments

JVRUST and JVCPP each preserve their ordering against the relevant SciPy baseline in all 201 non-batch cases. Required cases use native SciPy; finite-limit cases use capped SciPy where available. The original single-call suite retains square-expanded SciPy, so its wins must not be conflated with wins against capped SciPy. JVRUST's wins by suite are 11/12 single-call, 9/10 formulation, 43/47 scaling, and 112/132 extended cases in both runs. JVCPP has the same counts except 42/47 in scaling. All twenty extended losses are the same cases across runs.

JVRUST's congested base scaling crossover remains between $N = 64$ and $N = 128$; at 1024 × 1280 capped SciPy is 9.42 times faster than JVRUST in Python 3.14 and 9.15 times faster in Python 3.11. For

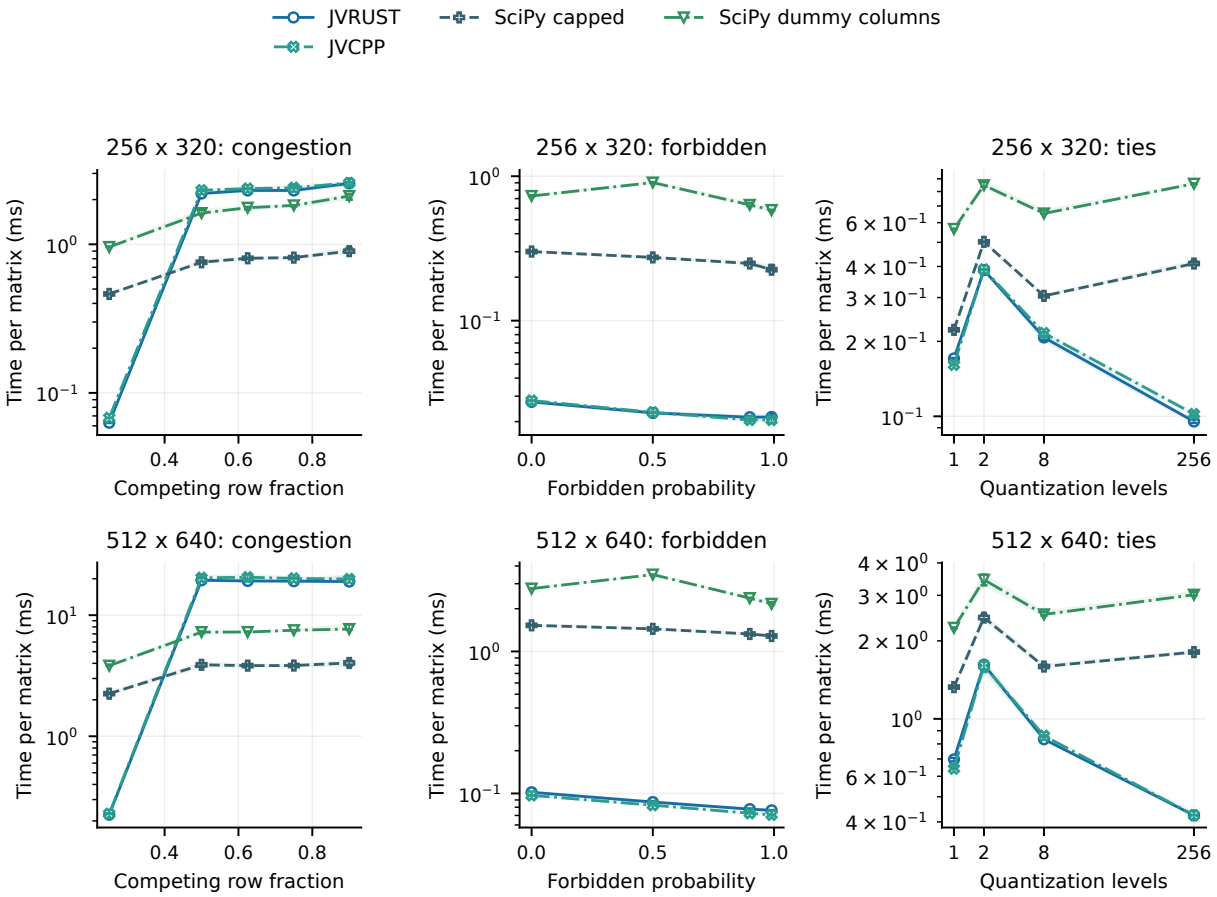


Figure 11: Finite-limit structure sweeps in Python 3.14. Limits are 0.5 for congestion, $2/M$ for forbidden entries, and 0.25 for quantized costs. Capped and dummy-column SciPy are included in every panel. Medians and timing-round interquartile ranges are shown.

1024 problems of size 8×10, JVRUST's loop-to-batch speedups are 2.39 and 2.30 with one worker, and 5.10 and 4.73 with four. Appendix A compares the main ratios, congestion, and ranking exceptions.

Consistency does not imply identical rankings. Across all solvers, 62 of 2,104 within-case pairwise median orderings reverse, and the fastest median changes in 31 of 204 cases. Requiring a margin exceeding 2% in both environments leaves twelve reversals; requiring 5% or 10% leaves one. These are post-hoc descriptive filters, not statistical significance thresholds. The 32×32 required case changes JVC/SciPy ordering: 10.87/9.71 microseconds in Python 3.14 versus 10.61/12.84 in Python 3.11. JVRUST remains faster than both.

The four-worker batch ranking also needs qualification. For 64 problems of size 64×80, JVCPP is approximately 37% slower than JVRUST in Python 3.14, but its 0.286 ms median is approximately tied with JVRUST's 0.293 ms in Python 3.11. The wide Python 3.14 JVCPP interquartile range weakens any claim of a stable implementation gap. The batching benefit itself repeats.

Median time(3.11)/time(3.14) ratios across equally weighted method/case settings are 0.915 for single-call, 0.913 for batch, 0.928 for formulation, 0.989 for scaling, and 1.011 for extended cases. These summaries describe the chosen grid, not an application mix or a universal speedup. Dependencies and session conditions differ, and additional independent sessions would be needed to assess small gaps or separate interpreter effects.

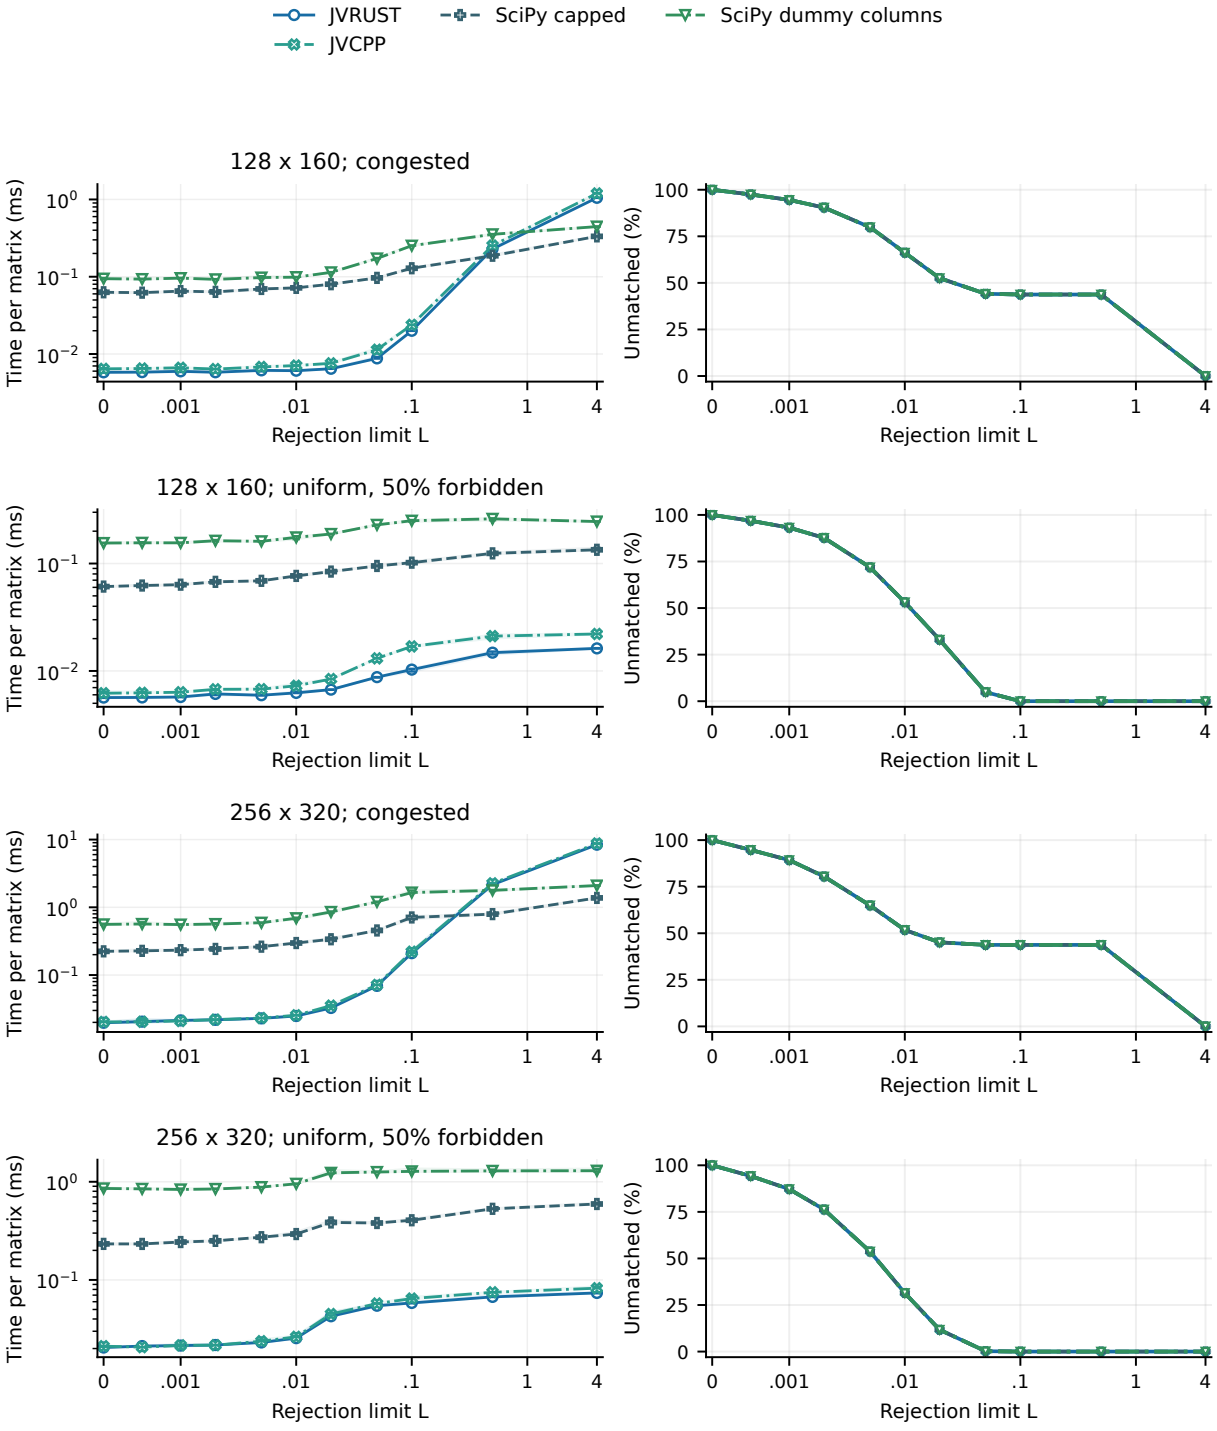


Figure 12: Rejection sweeps in Python 3.14 with the same input values reused across limits. Each row pairs runtime medians and timing-round interquartile ranges (left) with realized unmatched fractions (right), for one shape and cost family. The realized profiles agree exactly between environments for each recorded method and input.

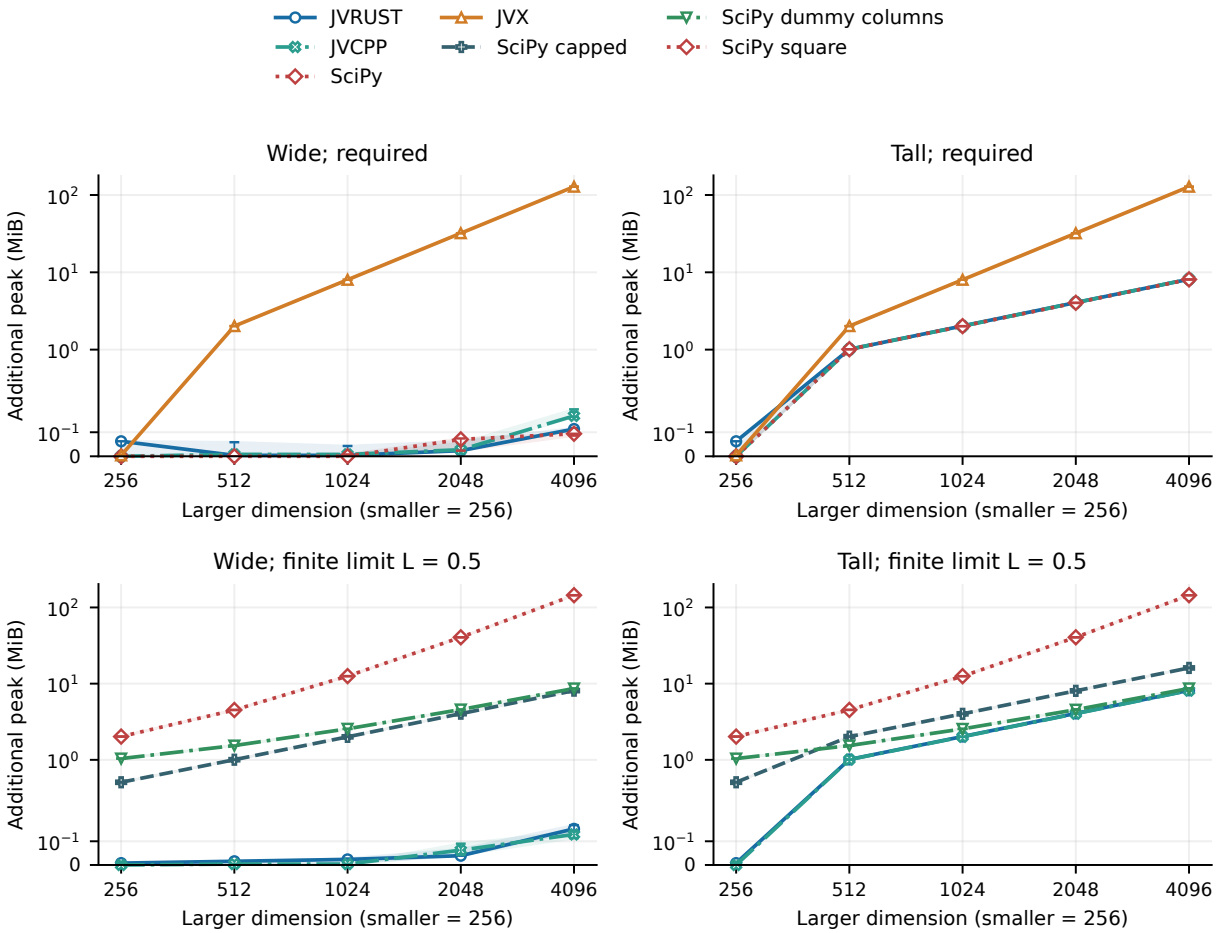


Figure 13: Python 3.14 additional peak working set for required and finite-limit assignment in both orientations. Each condition has five fresh processes; medians and interquartile ranges are shown. The scale is linear below 0.25 MiB to include zero increments. The input is already resident; these are high-water changes rather than exact allocation counts.

7 Discussion

7.1 Contribution relative to existing methods

The practical contribution is a uniform-penalty implementation with compatible Python calls, native rectangular storage, and reusable batch workers. Its rejection terminal stores one distance and one predecessor row. The experiments characterize these complete implementations; they do not isolate the gains from the terminal, reductions, scanning policies, or output adapters individually.

Establishing novelty or causal performance gains for individual policies would require comparisons with prior VJ and LSAPE techniques and ablations against an otherwise identical implementation. Capped SciPy controls dimensional expansion but differs in initialization, search, temporary storage, and adapter work.

7.2 Performance is determined by formulation and structure

The wide and finite-limit results demonstrate that transformed problem size can dominate practical interface performance. However, SciPy's required rectangular results and the block regression show why "no padding" cannot explain every speed difference. Both SciPy and the proposed solver avoid square padding for required rectangles, yet their performance ordering changes substantially with structure.

The conservative reduction bound in Section 4 also cautions against inferring a better asymptotic algorithm from favorable timings. Reductions can be effective on common inputs while consuming significant work on

competitive cost blocks. The block workload motivates phase-level profiling and controlled reduction-policy ablations before making stronger algorithmic claims.

The scaling curves make the dependence on both size and structure explicit. Increasing only the larger dimension differs from increasing both dimensions, and the congested finite-limit sweep changes the ordering of JVRUST and capped SciPy as size grows. The curves characterize these measured families without establishing an asymptotic exponent or predicting behavior on arbitrary matrices.

7.3 Implementation language and completion kernel

The three implementations compare implementation choices under a shared algorithmic design. JVHYBRID changes only square completion and provides no consistent whole-call benefit on the evaluated required square workloads. Differences on the shared rectangular Rust path should be treated as timing variation, not algorithmic differences.

JVRUST is faster than JVCPP in all twelve single-matrix cases in Table 3, with the largest gap on the tall input. The scaling sweeps do not preserve that ordering at every point: JVCPP has lower medians on several uniform cases, including the largest finite-limit fixed-row rectangle. These results compare particular implementations and compiler settings; they do not isolate a programming-language effect. The additional large cases keep the two implementations within about 6.1%, whereas extreme tall inputs and four-worker batches have different patterns. The two Python environments are not a controlled compiler or language experiment.

7.4 Implications for applications

On the evaluated synthetic costs outside the congested-block family, an implementation of the proposed solver had the lowest median among the configured methods in the original suites; the extension retains the native advantage against its selected SciPy baselines. The focused comparison does not rank excluded solvers. For the finite-limit congested family, JVRUST led at $N = 16$, 32, and 64, while capped SciPy led at $N = 128$, 256, 512, and 1024; these sampled sizes do not define a general crossover threshold. SciPy also led on the required 256×320 block case. In the original Python 3.14 formulation and finite-limit scaling suites, capped SciPy was the fastest evaluated SciPy formulation and used the smallest transformed cost buffer. This ranking is not universal: dummy-column SciPy can be faster on extreme tall inputs and has a lower additional peak working set in the tall memory sweep. The memory difference reflects the additional original-size buffer described in Section 6.9; the reported timings do not isolate the runtime cost of this adapter difference. In the Python 3.11 original scaling sweep at 64×1024, dummy-column SciPy also has a lower median than capped SciPy (158.17 versus 164.74 microseconds), with overlapping timing-round interquartile ranges. Native batching reduced latency for the tested batches of independent problems that were already available together; a tracker with only one assignment available at a time cannot exploit inter-problem batching without changing how work is scheduled. JVCPP provides a C++-only implementation, with workload-dependent performance relative to JVRUST. Application-specific matrices should determine the final choice.

An assignment acceleration is not an end-to-end application acceleration of the same factor. Cost construction, motion prediction, feature extraction, detection, and bookkeeping can dominate a tracking pipeline. This paper does not measure tracking accuracy, identity switches, or detector-plus-tracker frame rate.

Finite-limit integration also requires agreement on ties and numeric precision. Matching decisions at or near a penalty boundary may differ across implementations while retaining an equal or nearly equal objective. Tests should distinguish objective equivalence from a requirement for identical selected pairs.

8 Limitations and threats to validity

The performance study uses one Windows/hardware configuration and two software environments, four generated inputs per timing case, and nine timing rounds. NumPy, SciPy and native wheels differ, so the replication cannot isolate the effect of Python versions. CPU frequency and background activity are not controlled. Round interquartile ranges measure repeated-timing variation on the chosen inputs, not uncertainty over all possible assignment problems. Compiler and toolchain effects were not controlled; alternative builds may change the relative behavior of the implementations.

All timed inputs are float64. The native precision modes and conversion paths are validated more broadly by tests but are not characterized by the runtime tables. Very large dynamic ranges, subnormal penalties, and arbitrary floating-point near-ties require separate numerical studies.

The main finite-limit baselines use explicit square augmentation where no equivalent public limit API exists. Section 6.5 adds smaller dummy-column and original-dimension capped formulations, including a uniform-penalty specialization of [7]. Direct implementations corresponding to [5–7] remain important missing performance baselines. SciPy through our adapter is not a substitute for evaluating all specialized LSAPE kernels, and adapter work is part of the reported interface times.

The workload suite is synthetic. Several main finite-limit cases reject no smaller-side rows at the optimum; the formulation experiment covers rejection-heavy inputs and a congested finite-limit case. The scaling experiment covers fixed aspect ratios and a fixed smaller dimension, but uses only four random matrices per size and no held-out application distribution. Its linear and log-log curves are descriptive measurements, not empirical proofs of asymptotic bounds. These deliberately selected cases broaden coverage but do not establish application representativeness. Captured application matrices and held-out structured instances would improve external validity.

The original memory observations use one trial per condition; the added sweep uses five fresh processes and paired inputs per condition. They corroborate formulation and orientation effects but do not establish exact allocation counts, cross-platform resident-memory bounds, or long-running pool-retention behavior.

The three mathematical propositions justify the reformulations and final certificate in exact arithmetic. They do not prove all code paths, error cases, and floating-point transitions. Passing tests and exhaustive small cases complement, rather than replace, that distinction.

9 Conclusion

Thresholded rectangular assignment does not need a square matrix. A one-sided reformulation of the uniform penalty lets a Jonker–Volgenant search represent every rejection by one implicit terminal, and a final primal-dual certificate gives a clear criterion for correctness. The resulting solver, in Rust, hybrid, and C++ implementations, accepts the arguments of the expanded threshold interface without changing its objective.

The proposed implementations substantially outperform the expanded interfaces on the evaluated uniform and rejection-heavy workloads. JVRUST also outperforms capped SciPy in nine of the ten finite-limit formulation cases, although capped SciPy is faster on congested blocks and its advantage grows in the congested scaling family. Native batching reduces whole-batch latency on the tested independent workloads. The hybrid completion kernel provides no consistent benefit on the evaluated required square cases, and the Rust-versus-C++ ordering varies across scaling points. These comparisons establish results for the evaluated implementations, not superiority over every specialized optional-assignment solver.

The two-environment replication preserves the main native-versus-SciPy orderings across all tested sizes and structures, while revealing unstable small implementation gaps and four-worker batch rankings. Repeated memory trials also distinguish wide from tall native footprints. The contribution is a specific solver and a reproducible characterization within an established algorithmic framework. Further work should compare specialized LSAPE solvers, isolate reduction and initialization policies, explain the congested-block crossover, extend hardware and compiler coverage, and evaluate captured application workloads.

Data and software availability

The solver is distributed as laprust 1.0.0 [13] (JVRUST), lapcpp 1.0.0 [15] (JVCPP), and LAPX 0.11.1 [14] (JVRUST and JVHYBRID alongside the baseline families). The paper repository is github.com/rathaROG/lap-paper. The accompanying reproducibility artifact contains both environments' raw timing samples, memory trials, input hashes, schedules, validation records, package and binary identities, protocols, exact measurement sources, and figures. It also includes the cross-environment comparison, publication evidence checks, and arXiv build and HTML validation reports. All 17,892 timing samples and 830 memory observations are retained; the manuscript does not pool the environments into one timing estimate. Figures can be rebuilt from saved measurements without executing solvers.

Acknowledgements

The implementations retain attribution to Tomas Kazmar's lap, Vadim Markovtsev's src-d/lapjv, and Christoph Heindl and contributors to py-lapsolver. The assignment and error-correction literature provides the theoretical background and comparison points for this implementation study.

AI assistance disclosure

AI coding assistants (OpenAI Codex with GPT-6 Astra and GPT-6.1 Sol, and Anthropic Claude Code with Claude Opus 5.5) assisted with literature retrieval, source inspection, implementation changes, experiment-script preparation and execution, figure generation, and manuscript drafting. The raw artifacts identify the computations supporting the reported results. The author has cross-checked and reviewed the scientific claims, attribution, computations, and final text, and takes responsibility for the manuscript.

References

[1] Yifu Zhang, Peize Sun, Yi Jiang, Dongdong Yu, Fucheng Weng, Zehuan Yuan, Ping Luo, Wenyu Liu, and Xinggang Wang. ByteTrack: Multi-object tracking by associating every detection box. In *European Conference on Computer Vision (ECCV)*, 2022. URL https://arxiv.org/abs/2110.06864.

[2] Tomas Kazmar et al. lap. Source implementation. URL https://github.com/gatagat/lap. Evaluated through its LAPX adaptation.

[3] R. Jonker and A. Volgenant. A shortest augmenting path algorithm for dense and sparse linear assignment problems. *Computing*, 38:325–340, 1987. doi: 10.1007/BF02278710.

[4] David F. Crouse. On implementing 2d rectangular assignment algorithms. *IEEE Transactions on Aerospace and Electronic Systems*, 52(4):1679–1696, 2016. doi: 10.1109/TAES.2016.140952.

[5] Sébastien Bougleux and Luc Brun. Linear sum assignment with edition. Technical report, Normandie Université, 2016. URL https://arxiv.org/abs/1603.04380.

[6] William Jones, Aziem Chawdhary, and Andy King. Optimising the Volgenant–Jonker algorithm for approximating graph edit distance. *Pattern Recognition Letters*, 87:47–54, 2017. doi: 10.1016/j.patrec.2016.07.024. URL https://kar.kent.ac.uk/56765/1/main.pdf.

[7] Sébastien Bougleux, Benoit Gaüzère, David B. Blumenthal, and Luc Brun. Fast linear sum assignment with error-correction and no cost constraints. *Pattern Recognition Letters*, 134:37–45, 2020. doi: 10.1016/j.patrec.2018.03.032. First published online in 2018.

[8] J. Bijsterbosch and A. Volgenant. Solving the rectangular assignment problem and applications. *Annals of Operations Research*, 181:443–462, 2010. doi: 10.1007/s10479-010-0757-3.

[9] A. Volgenant. Linear and semi-assignment problems: A core oriented approach. *Computers & Operations Research*, 23(10):917–932, 1996. doi: 10.1016/0305-0548(96)00010-X.

[10] David F. Crouse. The Tracker Component Library: Free routines for rapid prototyping. *IEEE Aerospace and Electronic Systems Magazine*, 32(5):18–27, 2017. Routines `assign2DMissedDetect` and `assign2DFull`, https://github.com/USNavalResearchLaboratory/TrackerComponentLibrary, accessed 22 October 2026.

[11] The MathWorks, Inc. assignjv, assignmunkres, and assignauction. Sensor Fusion and Tracking Toolbox documentation, introduced in R2018b. URL https://www.mathworks.com/help/fusion/ref/assignjv.html. Accessed 22 October 2026.

[12] Vadim Markovtsev et al. lapjv. Source implementation. URL https://github.com/src-d/lapjv. Accessed 2026-10-22.

[13] Ratha SIV. laprust: Native Rust linear assignment solvers with implicit rejection and batch solving. Version 1.0.0, 2026. URL https://github.com/rathaROG/laprust.

[14] Ratha SIV. LAPX: Linear assignment problem solvers. Version 0.11.1, 2026. URL https://github.com/rathaROG/lapx.

[15] Ratha SIV. lapcpp: Native C++ linear assignment solvers with implicit rejection and batch solving. Version 1.0.0, 2026. URL https://github.com/rathaROG/lapcpp.

[16] Christoph Heindl et al. py-lapsolver. Source implementation. URL https://github.com/cheind/py-lapsolver. Evaluated through its LAPX adaptation.

[17] SciPy developers. Documentation for scipy.optimize.linear_sum_assignment. URL https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.linear_sum_assignment.html. Accessed 2026-10-22; experiments use SciPy 1.18.1 and 1.17.1.

[18] Russ J. Vander Wiel and Nikolaos V. Sahinidis. The assignment problem with external interactions. *Networks*, 30 (3):171–185, 1997. doi: 10.1002/(SICI)1097-0037(199710)30:3<171::AID-NET3>3.0.CO;2-I.

[19] Sébastien Bougleux, Luc Brun, et al. LSAPE bipartite matching toolbox. Version 4, 2017. URL https://bougleux.users.greyc.fr/lsape/.

A Replication with Python 3.11.17

The full base and extended collections were repeated with Python 3.11.17, NumPy 2.4.6, and SciPy 1.17.1 on 24 October 2026. Solver releases, host, input hashes, schedules, and measurement protocol match the Python 3.14.9 collection; dependencies and sessions differ. Tables A1–A3 provide the replicated single-call, batch, and formulation medians, and Table A4 compares within-environment ratios.

All correctness validations passed: 16,800 base enumeration calls including 80 expected infeasible calls, 10,520 formulation validation calls, 1,208 scaling checks, 2,112 extended timed-input checks, 129 edge checks, and the same package test counts as Section 6.1. The four-worker JVCPP and JVRUST medians on the 64×80 batch are approximately tied in Python 3.11, unlike the larger gap in Python 3.14. This appendix complements the full 201-case ordering comparison and repeated-memory findings in Sections 6.9–6.10.

A.1 Single-matrix runtime

Table A1: Median runtime in microseconds in the Python 3.11.17 run, with the cases and methods of Table 3.

Case	JVRUST	JVHYBRID	JVCPP	JVX	JVS	JVC	SciPy
Square 128×128	59.71	61.22	71.42	111.37	109.65	260.29	268.04
Square 512×512	1033.70	1047.44	1111.29	3265.40	3276.91	6031.73	5066.82
Wide 64×512	11.26	11.31	13.80	866.98	815.35	1803.09	76.49
Tall 512×64	23.76	23.28	78.40	887.86	832.86	1945.91	148.59
Rectangle 256×320	72.11	71.52	81.86	171.94	173.29	4014.82	559.79
Blocks 256×320	8991.83	8744.50	9801.58	9417.83	9580.05	5601.48	1443.51
Ties 128×160	49.30	53.29	52.25	59.57	60.04	93.52	77.25
128×160, $L = 0.5$	13.45	14.20	19.29	517.50	437.54	2178.79	454.00
256×320, $L = 0.5$	68.50	69.23	75.29	3652.79	3047.66	14849.92	2284.79
Tight 128×160, $L = 0.05$	9.17	9.43	14.28	451.08	378.58	2540.20	1083.24
Signed 128×160, $L = -0.25$	13.61	14.33	19.23	1021.79	848.29	3028.26	435.84
Square 128, $L = 0.5$	86.70	88.23	96.25	194.29	192.23	590.52	516.45

A.2 Batch runtime

Table A2: Median milliseconds per complete batch in the Python 3.11.17 run, with the workloads and methods of Table 4.

Method	$64 \times (64 \times 80)$ required, ms	$16 \times (128 \times 160)$ $L = 0.5$, ms	$1024 \times (8 \times 10)$ required, ms
JVRUST sequential loop	0.728	0.406	1.158
JVRUST batch, 1 worker	0.655	0.390	0.504
JVRUST batch, 4 workers	0.293	0.214	0.245
JVHYBRID batch, 1 worker	0.658	0.396	0.508
JVHYBRID batch, 4 workers	0.298	0.219	0.244
JVCPP batch, 1 worker	0.710	0.439	0.508
JVCPP batch, 4 workers	0.286	0.241	0.282
JVX batch, 1 worker	1.586	8.430	5.873
JVX batch, 4 workers	2.931	5.372	19.281
SciPy sequential loop	2.707	7.677	2.803

A.3 Formulation comparison

Table A3: Finite-limit formulation comparison in the Python 3.11.17 run, median microseconds per matrix, with the cases and methods of Table 6.

Case	Unmatched	JVRUST	JVCPP	Square	Dummy cols	Capped
128×160, $L = 0.5$	0.00	13.16	18.68	439.99	261.16	131.41
256×320, $L = 0.5$	0.00	71.34	77.14	2327.70	1356.99	577.33
128×160, $L = 0.05$	8.25	9.25	14.08	1054.59	256.30	103.35
Signed 128×160, $L = -0.25$	0.00	13.77	19.31	436.12	240.74	140.81
Square 128, $L = 0.5$	0.00	85.93	96.59	533.00	415.92	224.62
128×160, $L = 0.005$	87.75	5.83	6.87	825.64	170.22	73.17
128×160, $L = 0.001$	117.00	5.54	6.22	505.93	163.95	68.06
Tall 160×128, $L = 0.005$	89.50	13.47	12.05	880.12	175.97	84.32
Wide 64×512, $L = 0.005$	18.00	8.55	9.05	2072.51	197.94	89.00
Blocks 256×320, $L = 0.5$	112.00	1673.29	1723.54	2596.48	1635.13	729.36

A.4 Principal ratios in both runs

Table A4: Ratios of within-run median times, with numerator and denominator stated explicitly; t denotes median runtime. The last row is an absolute time. Each run's summaries are processed by the same calculation.

Quantity	Python 3.11.17	Python 3.14.9
t(JVRUST) / t(SciPy), required blocks (×)	6.23	6.20
t(capped SciPy) / t(JVRUST), smallest of nine wins (×)	2.61	2.50
t(capped SciPy) / t(JVRUST), largest of nine wins (×)	12.56	12.56
t(JVRUST) / t(capped SciPy), blocks 256×320 (×)	2.29	2.41
t(JVRUST) / t(capped SciPy), blocks 1024×1280 (×)	9.15	9.42
t(loop) / t(batch), $1024 \times (8 \times 10)$, one worker (×)	2.30	2.39
t(loop) / t(batch), $1024 \times (8 \times 10)$, four workers (×)	4.73	5.10
t(JVCPP) / t(JVRUST), largest finite-limit gap in Tables 3/A1 (×)	1.56	1.56
t(JVCPP) / t(JVRUST), tall 512×64 (×)	3.30	3.32
JVRUST, 128×160 at $L = 0.5$ in Tables 6/A3 (µs)	13.16	15.09

A.5 Consistent orderings and important exceptions

In both runs, JVRUST is faster than capped SciPy at $N = 16$, 32, and 64, while capped SciPy is faster at $N = 128$, 256, 512, and 1024. Figure A1 shows the unchanged native-versus-SciPy orderings and the sampled congested crossover. Points reflect the designed case grid, not independent draws from an application population.

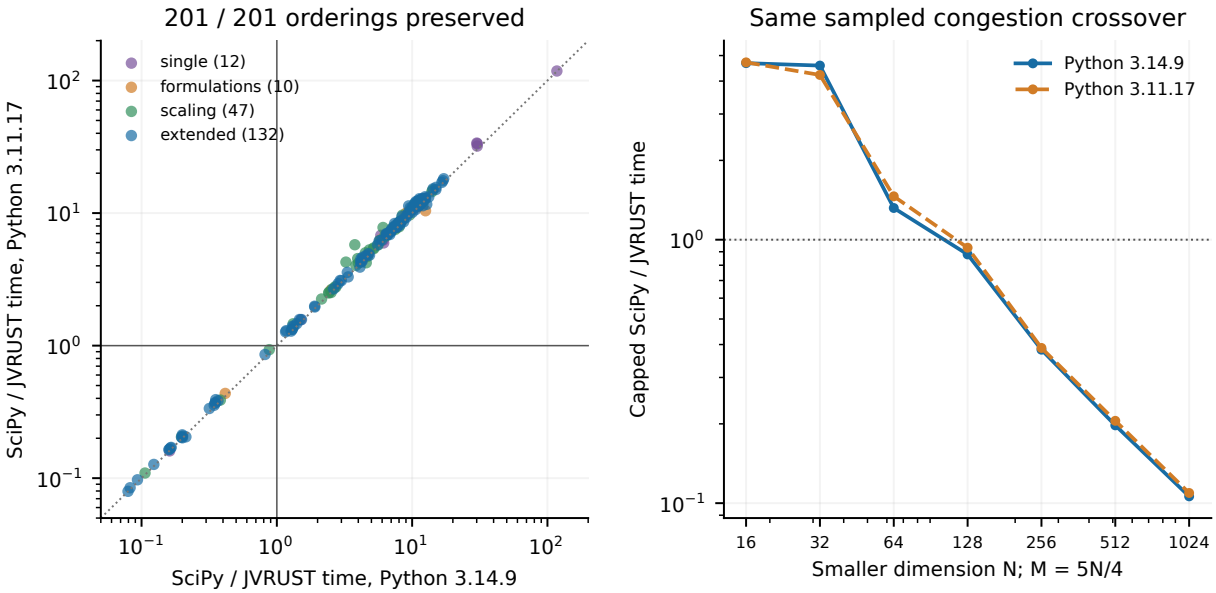


Figure A1: Ratios of within-run median times. Left: the relevant SciPy time divided by JVRUST time for each of the 201 non-batch cases in both environments, grouped by suite. Right: capped SciPy/JVRUST ratios in the congested base scaling sweep. Ratios above one favor JVRUST. Finite-limit single-call cases retain square-expanded SciPy; other finite-limit cases use capped SciPy.

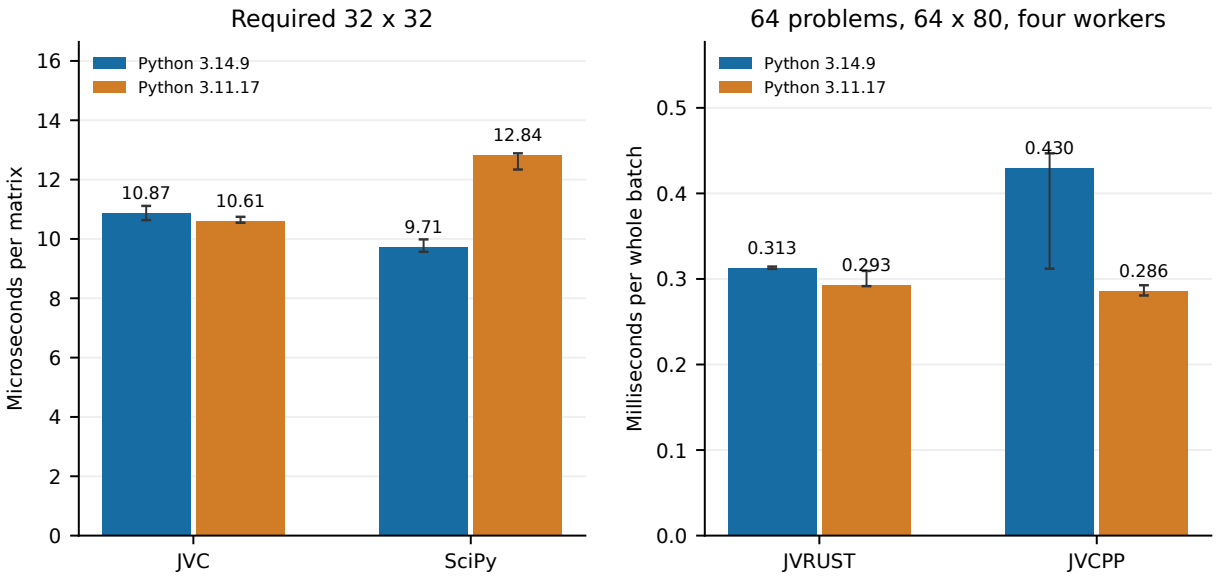


Figure A2: Two changes that qualify the consistency finding. Left: required 32×32 scaling, where JVC and SciPy exchange order. Right: 64 problems of size 64×80 with four workers, where JVCPP's large Python 3.14 gap becomes an approximate tie in Python 3.11. Whiskers are timing-round interquartile ranges, not confidence intervals. These examples do not overturn the repeated native-versus-SciPy ordering or the general batching benefit.